

Introduction to statistics

Wenda Zhou

June 25, 2017

Contents

1	Introduction	7
1.1	What is statistics?	7
1.2	Data	8
1.2.1	Rectangular data	8
1.2.2	Data types	8
1.3	Descriptive statistics	11
1.3.1	Measures of centrality	11
1.3.2	Measures of dispersion	11
1.3.3	Measures of association	11
1.3.4	Order statistics	12
1.3.5	Descriptive statistics for categorical data	12
1.3.6	Perils of descriptive statistics	13
1.4	Visualizing data	15
1.4.1	Visualizing one categorical variable	15
1.4.2	Visualizing one numeric variable	15
1.4.3	Visualizing two numeric variables	17
1.4.4	Visualizing one numeric and one categorical variable	18
1.4.5	Visualizing more than two variables	20
1.4.6	Visualizing other types of data	22
2	Probability	25
2.1	Probability axioms	25
2.1.1	Sample space and events	25
2.2	Calculus of probability	25
2.2.1	Disjoint or mutually exclusive events	26
2.2.2	Complement of an event	26
2.2.3	Inclusion-Exclusion	27
2.2.4	Independent events and multiplicative rule	27
2.2.5	Conditional probability	28
2.2.6	Conditional probability and independence	28
2.2.7	Conditional probability and additivity	29
2.2.8	Law of total probability	30
2.2.9	Conditional probabilities and bayes rule	30
2.3	Random variables	31
2.3.1	What is a random variable?	32
2.3.2	Random variables and events	32
2.3.3	Discrete random variables	32

2.3.4	Expectation	33
2.3.5	Population statistics	34
2.3.6	Continuous random variables	35
2.3.7	Distribution and density functions	36
2.3.8	Expectation (bis)	37
2.4	Common distributions	38
2.4.1	Bernoulli distribution	38
2.4.2	Binomial distribution	39
2.4.3	Poisson distribution	39
2.4.4	Uniform distribution	39
2.4.5	Normal distribution	40
2.4.6	Exponential distribution	40
2.5	Jointly distributed variables	40
2.5.1	Joint distribution	41
2.5.2	Joint distribution and events	41
2.5.3	Joint distribution and expectations	42
2.5.4	Marginal distribution	43
2.5.5	Independent random variables	43
2.5.6	Conditional distributions	44
2.6	Operation on random variables	44
2.6.1	Transforming a random variable	44
2.6.2	Sums of random variables	45
2.7	Properties of the expectation	46
2.7.1	Expectation of a sum	46
2.7.2	Expectation of product	47
2.7.3	Variance of a sum	47
2.8	Limit theorems	48
2.8.1	Law of large numbers	48
2.8.2	Central limit theorem	49
3	Sampling	51
3.1	Experiments	51
3.1.1	Randomized experiment	51
3.1.2	Control groups and placebo	51
3.1.3	Sub-populations and inductive inference	52
3.2	Observational studies	52
3.2.1	Prospective and retrospective studies	52
3.2.2	Confounding	53
3.2.3	Natural experiments	53
4	Estimation	55
4.1	Models and likelihood	55
4.2	Estimators	56
4.2.1	Unbiased estimators	56

4.2.2	Mean-squared error	57
4.3	Maximum likelihood estimation	57
4.3.1	Example: binomial model	57
4.3.2	Example: exponential model	58
4.3.3	Theoretical properties for the mle	58
4.4	Method of moments	59
4.4.1	Example: exponential distribution	59
4.4.2	Example: gamma distribution	60
4.5	Uncertainty in estimation	60
4.5.1	Confidence intervals	60
4.5.2	Confidence interval for a normal observation	61
4.5.3	Bootstrapping confidence intervals	61
5	Regression	59
5.1	Linear regression	59
5.1.1	Ordinary least squares	59
5.1.2	Residuals	61
5.1.3	Generalized linear models and link functions	61
5.2	Non-linear regression	62
6	Hypothesis testing	63
6.1	Vocabulary of testing and first example	63
6.2	Permutation tests	64
6.3	Pivot statistics	65
6.4	Testing contingency tables	65
7	Linear models	67
7.1	Coefficients of a linear model	67
7.1.1	Continuous variables in linear models	67
7.1.2	Categorical variables in linear models	67
7.1.3	Generalized linear models	68
7.2	Significance in a linear model	69
7.2.1	Testing significance of a single coefficient	69
7.2.2	Confidence interval for a single coefficient	70
7.3	Model selection	70
7.3.1	Information criterion	70
7.3.2	Which models to select from?	71
7.4	Penalized linear models	71
7.4.1	Ridge regression	72
7.4.2	Lasso regression	72
8	Bayesian statistics	75
8.1	Priors	75
8.2	Likelihood and posterior	75

8.3	Bayesian estimators	76
8.4	Examples	76
8.4.1	Binomial model	76
8.4.2	Normal model	77
9	Prediction	79
9.1	Defining prediction	79
9.1.1	What is a good prediction	79
9.1.2	Overfitting and underfitting	80
9.2	Linear models for prediction	80
9.2.1	Confounding in prediction	80
9.2.2	Uncertainty in prediction	81
9.3	Classification	81
9.3.1	Performance of classification	81
9.3.2	Linear models for classification	82
9.3.3	Linear discriminant analysis	83
9.3.4	Naive Bayes	84
10	Machine learning	85
10.1	Tree-based methods	85
10.1.1	Decision trees	85
10.1.2	Regression trees	85
10.1.3	Learning trees	87
10.1.4	Random forests	88
10.2	k -nearest neighbours	88
11	Advanced topics in statistics	91
11.1	Time Series	91
11.1.1	Stationarity	91
11.1.2	Measuring self-correlation	93
11.1.3	Common time series model	93
11.1.4	Seasonality	94
11.2	Mixed effects	94
11.2.1	ANOVA	95
11.2.2	Mixed effects	96
11.3	Missing data	98
11.3.1	Types of missingness	98
11.3.2	Non-response and response biases	99
11.3.3	Complete case analysis	99
11.3.4	Multiple imputation	100

1 Introduction

1.1 What is statistics?

The Merriam-Webster dictionary defines statistics as a branch of mathematics dealing with the collection, analysis, interpretation and presentation of masses of numerical data. In particular, we will be most interested in understanding how to work with and analyze data in the presence of uncertainty and unknowns. Indeed, it is most often the case that the problems we encounter have inherent variability (e.g. noise in the measurements) and that we do not know or understand the complete process being studied. It is thus important to still be able to derive conclusions we can be confident about despite the uncertainty.

The questions we will be interested in statistics can be roughly categorised into three broad types: prediction, estimation and testing.

Prediction Prediction is concerned about predicting quantities that are unknown to us by making use of related information that is known to us. For example, a weather forecaster predicts the weather tomorrow using meteorological measures today and past information. Amazon predicts which products its customers are most likely to buy.

The problem of prediction has enjoyed an important revival in the past decade with the increase in computing power and the increasing amount of data being collected throughout the world, and is the central problem of machine learning.

Estimation Estimation is concerned about assigning values and uncertainty to quantities that are unknown and most often cannot be observed directly. For example, an economist may be interested in estimating the average effect (for example in lifetime earning increase) of obtaining a university degree. A pharmacologist may be interested in estimating the average effect (for example, in increase life expectancy) of a cancer treatment.

The problem of estimation is historically the central problem that motivated statistics. Although it shares many common aspects with prediction, it also differs in subtle but important ways.

Testing Testing is concerned about making decisions in the face of uncertainty, and is a problem that is closely related to that of estimation. For example, a drug company may be interested to understand whether a drug is more effective than the placebo. An advertiser may be interested in whether their advertisement is effective.

The problem of testing is very closely related to that of estimation, as it will often be the case that the test we wish to understand can be phrased in the sense of “is the effect

1 Introduction

zero”?. However, the language the notion of testing provides will prove valuable, and the fact that it is such a common problem warrants a separate mention.

1.2 Data

As the amount of data collected in the world increases, the diversity and variety of the data collected also increases. However, the vast majority of the data can still be understood in a simple rectangular fashion we describe below.

1.2.1 Rectangular data

It is often the case that we may think of a dataset as a collection of *observations* each having a collection of characteristics (often called *variables*). Most often, we are interested in how some variables (often called dependent or response variables) change or vary as a function of some other variables (often called *independent* or *explanatory* variables).

Example 1 (Clinical trial) Suppose a drug company did a clinical trial with a 100 patients for a drug designed to lower blood cholesterol. For each of them, they recorded the weight, age, sex, the blood cholesterol before and after taking the drug.

In this example, each patient corresponds to an observation, and the variables are the weight, age, sex, and blood cholesterol. In the context of a drug trial, we are interested in how effective the drug is, so the response variable could be the blood cholesterol after taking the drug, or maybe the difference in blood cholesterol before and after taking the drug. If we suspect that not everyone will respond similarly to the drug (e.g. the drug might be more effective for women than men), then we may be interested to consider the other variables as explanatory variables.

One interesting aspect in this case is that we may also want to consider the cholesterol before taking the drug as an explanatory variable. Indeed, if we believe that the effect of the drug depends on the initial amount of cholesterol (e.g. the drug works particularly well for people with very high levels of cholesterol, but not for others), then we certainly would want to understand it.

It will often be convenient to collect all the variables for all the observations into a *data matrix* or *data frame*. This is a rectangular table with each row corresponding to an observation and each column corresponding to a variable. It is usual to let n be the number of observations in the data frame.

Example 2 (Clinical trial (continued)) Suppose that we consider the same study as in example 1. We may collect all of the information into a data frame as below.

1.2.2 Data types

A given variable in a dataset will often have some restrictions on the values it can take, which we will refer to as the type of the variables. In example 1, the variable “sex” can only take two values (M / F), neither of which are numbers. On the other hand, the

Age	Sex	Weight (kg)	Cholesterol before (mg / dL)	Cholesterol after (mg / dL)
41	M	95	245	235
50	F	85	250	230
⋮	⋮	⋮	⋮	⋮

variable “weight” could (potentially) take any non-negative value. Most of the data we will study can be classified as either *numerical* or *categorical*.

Numerical variables

A numerical variable is a variable that represents a quantity and can take a range of numerical values. The quantity represented can be discrete such as a count (which may only take values 0, 1, ...), or continuous (e.g. the concentration of blood lipids). In addition, the range of a numerical variable may often be restricted: some quantities are restricted to be non-negative (such as weight), and other quantities may be restricted to be between 0 and 1 (for example, the proportion of patient experiencing a side effect).

Be careful that not all numbers are numerical variables! For example, phone number or a zip code is not a numerical variable. Indeed, although they are numbers, they do not represent a quantity and we cannot add, subtract or average them.

A somewhat special yet oft encountered numerical variable type, *circular* variables represent quantities with somewhat unusual arithmetic properties, as they are values that “wrap around”. Common examples of circular data include time of day (e.g. 12:59pm is intuitively “close” to 1:01pm) or day of the year. We should take special care when operating on such variables and computing quantities such as averages.

Categorical variables

A categorical variable is a variable that represents groups or categories. A categorical variable usually takes on a finite number of possibilities known in advance (for example, male / female, or one of the 50 states of the U.S.). Each such category is called a *level*. In addition, all these examples display no particular order between the levels, and so are said to be *nominal*.

On the other hand, some categories may have a natural ordering. For example, suppose we asked the following question on a survey: “how often do you go to church”, and offered the following possible answers:

1. Less than once a year,
2. A few times per year,
3. A few times per month,
4. Every week,
5. Every day,

1 Introduction

then it is clear that these categories have a natural ordering to them. We call such categories *ordinal*.

Likert scales

A very common type of variables in surveys and other similar datasets are the so called *Likert* scales (name after its inventor, psychologist Rensis Likert). The typical five-point likert item is of the form: “strongly disagree”, “disagree”, “neither agree nor disagree”, “agree”, “strongly agree”, although similar types of answers are also usually referred to as a Likert item.

From our previous discussion, Likert scales fall squarely into the ordinal categorical variable categories, with each category being ordered with respect to each other, but the variable not representing a specific numeric quantity. However, it is very common in practice to treat such scales as numeric, assigning for example a value from 1 to 5 to each item. Indeed, this simplification often works well in practice, especially when summarising large ensembles of Likert items.

Other types of data

Although the most common types of data fall within the numeric or categorical types, an increasing portion of the data collected today does not necessarily belong to either of those types, or displays subtle differences and will require special treatment. We mention a couple of such data types for completeness, although we will not have the opportunity to study them in this course.

Graphs and networks An increasingly important type of data today emanates from the relationship between different entities. For example, social networks present a rich structure by examining for example each user’s friend or contact list, and the declared interests of each user. Such data is usually best summarised into a *graph*, which is often used to capture relationships between various entities (e.g. users).

Images, sound and other signals In recent years, there has been huge progress in the analysis of images, sounds and other types of signal (scientific imaging, neural data etc.). Google announced in May 2017 that their image recognition now outperforms humans on a benchmark dataset. Although the simplest models simply treat such signals as high-dimensional numeric variables (e.g. treating each pixel of the image as a number), the best models attempt to make use of the specific structure of those signals.

Text Parallel to the improvement in our ability to learn and process signals, the last five years has seen rapid improvement in our ability to understand text, most notably in fields such as machine translation and sentiment analysis. Although text is inherently a discrete structure, it is neither categorical nor numeric. The best models attempt to learn an efficient numeric representation of text.

1.3 Descriptive statistics

Datasets are potentially complex objects with numerous variables, and so it is often useful for us to be able to synthesise the information of the entire dataset into some numbers, or descriptive statistics. These values that we compute are also called sample statistics, and hence will be referred to as *sample* (name-of-the-measure). In chapter 2, we will be instead looking at their theoretical counterparts, which will sometimes be called *population* (name-of-the-measure).

1.3.1 Measures of centrality

A *measure of centrality* is a number that attempts to summarise the location of the data in bulk, the two most common and well known being the (arithmetic) *mean* and the *median*.

The mean of a sample of n real-valued observations x_1, \dots, x_n is often written \hat{x} , and is defined as:

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i. \quad (1.1)$$

The median of a sample of n real-valued observations x_1, \dots, x_n is the value of the observation such that half of the observation are smaller, and half of the observations are larger. This notion can be generalized to the notion of a *percentile*. The p^{th} percentile is the value such that $p\%$ of the observations are *below* the given value. The median is then the 50^{th} percentile. The often used first and third *quartiles* can also be defined as the 25^{th} and 75^{th} quartile.

1.3.2 Measures of dispersion

A *measure of dispersion* is a number that attempts to summarise the spread of the data. The most common measures are the variance (and its cousin the standard deviation), and the interquartile range.

The variance of a sample of n real-valued observations x_1, \dots, x_n is often written σ^2 , and is defined as

$$\sigma^2 = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2, \quad (1.2)$$

where \bar{x} is the mean defined as in eq. (1.1). The standard deviation is then simply σ , the square root of the variance. The variance and standard deviation are also always non-negative quantities.

The *interquartile range*, often written *IQR*, is defined as the difference from the third quartile to the first quartile.

1.3.3 Measures of association

A *measure of association* is a number that attempts to summarise the association of two variables – i.e. how related they are. The most common such measure is the *covariance*,

1 Introduction

and its normalized version the *correlation*.

The covariance of two samples of n real-valued observations each, x_1, \dots, x_n , and y_1, \dots, y_n , is sometimes written σ_{xy} , and is defined by:

$$\sigma_{xy} = \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y}). \quad (1.3)$$

Note that we may interpret the variance of x_i as the covariance of x_i with itself. The correlation is a normalized version of the covariance, and is a unit-free quantity, defined by:

$$\text{corr}_{xy} = \frac{\sigma_{xy}}{\sigma_x \sigma_y}, \quad (1.4)$$

where σ_x, σ_y are the standard deviations of the x_i and y_i respectively. Note that the correlation is always between -1 and 1 , and has the same sign as the covariance. When the correlation is positive, we say that x and y are *positively correlated*. When the correlation is negative, we say that x and y are *negatively correlated*. When x and y are positively correlated, larger values of x tend to lead to larger values of y , whereas when x and y are negatively correlated, larger values of x tend to lead to smaller values of y .

1.3.4 Order statistics

For any sample of n real-valued observation x_1, \dots, x_n , we may be interested in the largest (or smallest) value. More generally, we may be interested in the p^{th} smallest (or largest) value. These values are called order statistics, and are usually written $x_{(p)}$ (note the parentheses, and pronounce “ x order p ”). By definition, $x_{(1)}$ is the first smallest – or simply smallest – value of the sample, whereas $x_{(n)}$ is the n^{th} smallest – or largest – value of the sample.

1.3.5 Descriptive statistics for categorical data

The descriptive statistics we have seen so far are inherently adapted to describing numerical data. However, they have no meaning when we consider categorical data. The most common summary for purely categorical data is called the *contingency table*, which collects the count of occurrences of each category or combination of categories.

Indeed, suppose that we collect information concerning the hair colour (blonde, red, brown or black) and eye colour (blue, green, brown or black) of 20 individuals. In the rectangular data format that we are used to, that would correspond to 20 observations of 2 variables each, as in table 1.1. However, as the order of the observations does not matter, a way of summarising the data is the two-way contingency table, which records the number of individuals for each combination of eye colour and hair colour, as in table 1.2

In addition, we may also choose to ignore one or the other characteristic. Suppose for example that we only look at eye colour, ignoring hair colour. That is, we count the number of people with the given eye colour no matter what their hair colour is. We

Observation #	Hair Colour	Eye Colour
1	Brown	Blue
2	Blonde	Brown
\vdots	\vdots	\vdots

Table 1.1: Sample of data collected from 20 individuals

Hair colour	Eye colour			
	Blue	Green	Brown	Black
Blonde	2	1	2	1
Red	1	1	2	0
Brown	1	0	4	2
Black	1	0	2	0

Table 1.2: contingency table of eye and hair colour for 20 individuals

Eye colour	Blue	Green	Brown	Black
	5	2	10	3

Table 1.3: Marginal distribution of eye colour

would then obtain a one-dimensional table, called the *marginal* table or distribution of hair colour. For example, see table 1.3.

On the other hand, instead of ignoring the hair colour, we may choose to only look at people with the given hair colour. This corresponds to looking at a single row of the two-way table, and is called the *conditional* table or distribution.

The notion of contingency table can be extended to more than two variables, but we are effectively *adding* a dimension for each variable. For example, suppose that we had also recorded the gender of the person in the previous example. We could then have a $4 \times 4 \times 2$ table of all the possibilities, but this can be difficult to present. A possibility is to present two 4×4 tables, one corresponding to men, the other to women. However, as the number of categories and variables increases, this unavoidably becomes more complex as the data itself becomes more complex.

1.3.6 Perils of descriptive statistics

Descriptive statistics are a convenient way of summarising often complex datasets. Due to their simplicity, they may however fail to capture the full complexity of the dataset. A common example is Anscombe's quartet, a collection of four samples plotted in fig. 1.1 that have the same mean and variance of both x and y , and the same covariance, but inherently different properties.

Other common features that descriptive statistics often fail to capture are for example




Figure 1.1: Anscombe's quartet, four datasets with the same mean, variance and covariance. Source: Schutz

multi-modal datasets, where the data can be divided in several groups, or data with a lot of structure in the outliers. It is therefore often advantageous to not only look at summary statistics but also graphical representations of the data to better understand its structure.

1.4 Visualizing data

The first step in any analysis of the data is understanding the large structures that may exist in the data. The best way to achieve this is through graphing and visualizing data. By making use of the appropriate visualization, we will be able to understand better the data collected for a single variable, or how two (or more) variables relate to each other. As variables have different types, we will need different methods to visualize such variables.

In this section, we will be illustrating the techniques using a dataset of 1035 records of heights and weights of MLB players obtained from the SOCR. For illustration, we have included the first 10 rows of the dataset in table 1.4. We note that the team and position variables are categorical, whereas the height, weight and age variables are numeric.

1.4.1 Visualizing one categorical variable

A single categorical variable can be summarised by simply the counts (or proportions) of each of its categories. Suppose in our example that we wish to understand if some positions are more represented than others in the dataset. The usual method for displaying such results is the *bar chart*, illustrated in fig. 1.2. The bar chart aggregates each categories by the number of responses in the given category, and plots each category side by side.

What not to do! Another unfortunately common visualization for such types of data is the infamous *pie chart*, as in fig. 1.3. However, it suffers from many problems, due to the fact that humans have difficulty comparing areas. For example, looking at fig. 1.3, it is difficult to compare the relative size of second basemen and shortstops, and a question such as whether there are twice as many relief pitchers as outfielders is extremely difficult to answer. Whenever you feel a pie chart would be an adequate visualization of the data, a bar chart is nearly always more appropriate.

1.4.2 Visualizing one numeric variable

Histograms Suppose we are now interested in visualizing the heights of players in the MLB. A common technique to visualize one numeric variable is the *histogram*, as seen in fig. 1.4. Such a graphic presents the count (or sometimes the proportion) of players having the height in the given bin. From the histogram, it is for example easy to see that nearly all players have a height between 70 and 80 inches. However, note that histograms can be very sensitive to the bin width, especially if the data is discrete. It is often a good idea to try a few different widths.

1 Introduction

Table 1.4: First 10 row of SOCR MLB player data set

Name	Team	Position	Height (in)	Weight (lbs)	Age (yr)
Adam Donachie	BAL	Catcher	74	180	22.99
Paul Bako	BAL	Catcher	74	215	34.69
Ramon Hernandez	BAL	Catcher	72	210	30.78
Kevin Millar	BAL	First Baseman	72	210	35.43
Chris Gomez	BAL	First Baseman	73	188	35.71
Brian Roberts	BAL	Second Baseman	69	176	29.39
Miguel Tejada	BAL	Shortstop	69	209	30.77
Melvin Mora	BAL	Third Baseman	71	200	35.07
Aubrey Huff	BAL	Third Baseman	76	231	30.19
Adam Stern	BAL	Outfielder	71	180	27.05




Figure 1.2: Bar chart of position




Figure 1.4: Histogram of heights

Density estimates In addition to the sensitivity to the bin width, histograms have an unfortunate characteristic of being quite sensitive to the placement of the bin edges. Kernel density estimators (often abbreviated *kde*) are an alternative way of displaying the same data without defining specific breakpoints. They can also be seen as estimators for a distribution's density, which we will discuss in the probability section. An example of a kde is given in fig. 1.5. KDEs feature a parameter similar to the bin width of a histogram, usually called the *bandwidth*.




Figure 1.3: Pie chart of position

1.4.3 Visualizing two numeric variables

In this section and the next, we will be interested in visualizing how two variables relate. First, suppose that we wish to consider how the weight of a player is related to its height. In order to do so, we will use the *scatter plot*. The scatter plot (see fig. 1.6) displays one point for each observation (in this case each player), with the coordinate of the point determined by the two variables under consideration (in this case, height and weight).




Figure 1.5: Kernel density estimate of heights

In this case, the scatter plot displays clearly the increasing relationship of height with weight.

Scatter plots are particularly adequate to represent numeric variables when the variables themselves are not related to the observations. For example, in this particular case, each observation is a different player, who could (in principle) have any height or weight. However, another common scenario is when one of the variables indexes the observations, for example in the case of time series. In this case, a line plot is a good choice to display trends and patterns in the data across time. For example, fig. 1.7 displays the usage of the Capital bikeshare program in Washington D.C. during the first week of June 2011.

1.4.4 Visualizing one numeric and one categorical variable

Boxplot Returning to the baseball player dataset, suppose now that we wish to understand how the weight of the players differ according to the position. One way to do so would be to group the players according to their position, and then plot some summary statistics for each the groups. The most common such plot is the well-known *boxplot*, which represents the median, quartiles and outlying observations of the data.

The standard boxplot contains three main parts (see fig. 1.8): a middle box with a line, which represents the first and third quartile (with the middle line representing the second quartile or median), whiskers on either side of the box (representing some deviation), and outlying points. In R, the convention is for the whiskers to extend to furthest away whilst still within $1.5 \times \text{IQR}$. Points beyond that distance are then plotted individually.




Figure 1.6: Scatter plot of heights and weight




Figure 1.7: Bikeshare program usage




Figure 1.8: Anatomy of a boxplot

To visualize the relationship between a numeric and categorical variable, we may then draw one box representing each group, and put them side by side. For the example with the baseball players, see fig. 1.9.

Violin plots An alternative to the boxplot takes inspiration from the kernel density plot (see section 1.4.2) by attempting to draw one kernel density estimate for each category and arrange them vertically. This can be particularly useful when one suspects that the data is multi-modal and cannot be adequately described by its quartiles. On the other hand, the plot is significantly more complex, and computing reasonable density estimates requires more data points, hence boxplots may be preferable when each of the category is small.

1.4.5 Visualizing more than two variables

It is often more difficult to visualize the relationship between more than two variables as our visual system is particularly suited to recognise patterns in 2 dimensions. The main ideas to visualize more than two variables rely on either super-imposing several plots, or putting them side by side.

Faceting The idea of *faceting* plots is quite similar to that of conditional contingency tables. Instead of visualizing three variables simultaneously, we choose two variables and produce a plot using any of the techniques described above. However, instead of using all the observations in the plot, we instead split the observations according to the value of the third variable, and display those plots side by side.

For examples, in fig. 1.11, we have displayed side by side a boxplot for the number of users of the bike share program by hour of the day (where we are treating hour as a categorical variable). The plot is split between working days and non-working days, allowing us to easily visualize the different types of usage.

Other characteristics An alternative to creating several plots is to superimpose existing plots and represent the third variable by using a characteristic other than the position




Figure 1.9: Boxplot of weight by position




Figure 1.10: Violin plot of weight by position

1 Introduction




Figure 1.11: Usage of the bike sharing program by hour of day for working days and non-working days

of the graphic. For example, we may try to use a different colour between working and non-working days.

Different types of graphics will have different secondary characteristics that can be leveraged to represent further variables beyond the first two. For example, in addition to a points x and y coordinate, points in a scatter plot may also have different sizes and colours. In a line plot, we may use different colours and line types. In a bar chart, we may use different colours and fill patterns.

1.4.6 Visualizing other types of data

As we mentioned in section 1.2.2, there are numerous other types of data that are not easily represented in a simple rectangular format. Subsequently, these types of data are also difficult to visualize using the standard tools described above, and may require specially developed tools. For example, the problem of graph visualization is a rich field of its own right, with many existing tools. Similarly, the scientific communities have developed numerous visualization techniques adapted to the problems they are faced with.




Figure 1.12: Usage of the bike sharing program by hour of day for working days and non-working days

2 Probability

Probability is the language that enables us to discuss uncertainty in a mathematical fashion. The modern theory of probability is axiomatized by the so-called Kolmogorov axioms of probability, which attempt to encode our intuitive notion of probability.

2.1 Probability axioms

2.1.1 Sample space and events

The basic element of probability theory is an event. Philosophically, an event is a contingent proposition, that is, an assertion that may or may not be true. For example, one could consider events: “the coin lands on head”, “the die lands on an even number”, “it will rain tomorrow”. The probability axioms then describe the rules any quantification of uncertainty (i.e. probability) of events must obey.

It is mathematically convenient to use the language of set theory to describe events in the following fashion. We will use the letter Ω to refer to the set (or “collection”) of *all* possible outcomes. We call Ω the *sample space*.

Example 3 Sample space for a coin flip. There are only two possible outcomes for a coin flip: heads (write H) and tails (write T). We may then collect all these outcomes into a set $\Omega = \{H, T\}$, the set with the two elements H and T .

Sample space for a die roll. For a standard 6-sided die, there are 6 possible outcomes corresponding to the numbers 1 to 6. Hence we may represent the sample space as the set $\Omega = \{1, 2, 3, 4, 5, 6\}$.

Sample space for the weather. The sample space for the weather is somewhat too large to describe as explicitly as we have for the previous two examples. However, we will see that our inability to formulate it in a precise fashion does not impede our ability to talk about uncertainty in quantitative way.

Given a sample space, we may then consider events in that space. In the case of a coin flip, an event could be that the coin comes up heads. In the case of a die roll, an event could be that the die lands on 3. However, an event could also be that the die lands on a 1 or a 3. This suggests that an event is a collection of outcomes – or more mathematically, a *subset* of the sample space. We will thus write $E \subseteq \Omega$ for an event E .

2.2 Calculus of probability

In order to quantify uncertainty of events, we will assign to each event E a number between 0 and 1, the probability of the event, and write $P(E)$ for the probability of E .

2 Probability

One possible interpretation for this number is the long run proportion of events that are true: for example, we say that the probability of a fair coin landing on heads is 0.5 as if we would repeat the coin toss “infinitely” many times, half of them would land on heads – this is the so-called *frequentist* interpretation of probability. Another interpretation of this number is that it represents one’s personal beliefs about the likelihood of success of a particular event – for example, we would say that the probability of a fair coin landing on heads is 0.5 as we would be willing to take a 1 for 1 bet on the coin landing on heads. This is the so-called subjective or *Bayesian* interpretation of probability. Both points of view will be useful for the statistician, but we note that the calculus of probability – that is, the rules we use to manipulate probabilities, remain the same.

In addition to assigning a number between 0 and 1 to each event, we will have some rules that relates the probability of related event – for example, we would like to codify the intuition that no matter which die we use, the event “the die lands on an even number” is more likely than the event “the die lands on the number 2 exactly” as the former includes the latter.

First, let us mention two special events, \emptyset and Ω . Ω is simply the set of all possible outcomes, and we have that $P(\Omega) = 1$, which we may interpret as the fact that of all the possible outcomes, one must happen. Its counterpart, \emptyset , is the empty set of outcomes, and we have that $P(\emptyset) = 0$. We may also interpret these sets as events, in which case Ω is the event that always happens, whereas \emptyset is the event that nothing happens.

2.2.1 Disjoint or mutually exclusive events

We now come to the rules of the calculus of probability, which relates the probability of related events. Consider two events A and B , what can we say about the probability of $A \cup B$: the event that either A or B (or both) happen? In general, this is a difficult question, as A and B may overlap a little (or a lot). However, in the special case that A and B do not overlap, that is, it is impossible that both A and B happen, we have the simple rule of additivity:

$$P(A \cup B) = P(A) + P(B) \quad (2.1)$$

We say in this case that A and B are *disjoint* or *mutually exclusive*. We may also express the condition that A and B are mutually exclusive by the following mathematical assertion: $A \cap B = \emptyset$, which translates that there is no outcome where A and B both happen.

2.2.2 Complement of an event

Given an event A , we can consider the event that A does not happen. This event is called the *complement* of an event, and is written A^C . What can we say about the probability of A^C ? Well, first, note that either A or A^C happens (that is, something either happens or does not happen).

Mathematically, we may codify this statement into the following equality:

$$A \cup A^C = \Omega, \quad (2.2)$$

which states that A , together with its complement A^C , are equal to the entire sample space.

Now, by definition, we also have, that A and A^C are disjoint (specifically, A^C was defined in this way). We may codify this statement mathematically as

$$A \cap A^C = \emptyset. \quad (2.3)$$

Now, this means that we may apply our previous rule for additivity of disjoint events, and write $P(A \cup A^C) = P(A) + P(A^C)$. However, by using eq. (2.2), we see that $P(A \cup A^C) = 1$. Hence, we deduce the following relationship between the probability of A and its complement:

$$P(A) + P(A^C) = 1, \quad (2.4)$$

which can also be written as $P(A^C) = 1 - P(A)$.

2.2.3 Inclusion-Exclusion

In some cases, we may be interested in calculating $A \cap B$ despite A and B not being disjoint. In that case, the additive formula does not directly apply. However, we may apply it with a slight modification, giving the so-called *inclusion-exclusion* formula.

To understand the inclusion-exclusion formula, it is important to understand why it is not true that $P(A \cup B) = P(A) + P(B)$. By thinking for example about the case where $A = B$, we may see that it is due to the fact that we are counting the events that fall in both A and B twice on the right hand side.

In order to adjust for this, the inclusion formula has a correction on the right hand side, giving the formula:

$$P(A \cup B) = P(A) + P(B) - P(A \cap B). \quad (2.5)$$

Example 4 Suppose that we still consider a die roll, and consider the events $A = \{1, 2\}$, and $B = \{2, 3\}$. We have that $P(A) = P(B) = 1/3$, and note that $P(A \cup B) = 1/2$.

Now, the additive formula would claim $P(A \cup B) = 2/3$, which is clearly wrong. Let's use the inclusion-exclusion formula instead. To do so, we need to compute $P(A \cap B) = P(\{2\}) = 1/6$, to finally obtain:

$$P(A \cup B) = \frac{1}{3} + \frac{1}{3} - \frac{1}{6} = \frac{1}{2}, \quad (2.6)$$

as we expected.

In this case, we can see that without the correction term, we were counting the probability of the roll being 2 twice, making the calculation incorrect.

2.2.4 Independent events and multiplicative rule

Given an event A and an event B , we will often be interested in a very special relationship between the two events, called *independence*. Conceptually, independence denotes the fact that one event does not inform about the other. For example, consider a die roll, and

2 Probability

the following two events $A = \{2, 4, 6\}$ (i.e. the result is even), and the event $B = \{1, 2\}$ (i.e. the roll is 1 or 2). Then knowing that the result was 1 or 2 does not give us any information about whether the result was even or odd.

Mathematically, we say that A and B are independent if we have that

$$P(A \cap B) = P(A)P(B). \quad (2.7)$$

It is not immediately obvious how this mathematical statement relates to the fact that the event A does not give us any information about event B and vice versa, but we will explore this a bit more when defining conditional probabilities. If we apply the example to the A and B given above, we have that $P(A) = 1/2$, $P(B) = 1/3$, and $A \cap B = \{2\}$ thus $P(A \cap B) = 1/6$, and everything works as expected.

We will be revisiting independence later for random variables, as it is one of the most important concepts in probability and underpins much of statistics.

2.2.5 Conditional probability

Conditional probability allows us to reason about random events about which we have received partial information. For example, suppose again that we are throwing a die, and suppose that we would like to know the probability of the roll being 4 or higher (i.e. 4, 5, 6). A priori, this happens half of the time. However, suppose now that we roll the die, and do not disclose the result, but only that the roll was odd. How should this change the probability of the roll being 4 or higher? Intuitively, this seems to make our chance worse, as there is only one odd number in 4, 5, 6, but two even numbers.

Let us then define the condition probability of an event A (in this case $A = \{4, 5, 6\}$) given event B (in this case $B = \{1, 3, 5\}$, the event that the roll was odd). Write the conditional probability $P(A | B)$ (read “probability of A given B ”), and define:

$$P(A | B) = \frac{P(A \cap B)}{P(B)}. \quad (2.8)$$

Note that as we are dividing by the probability of B , we require B to have positive probability (i.e. $P(B) > 0$).

Example 5 Let’s compute out our example in detail. We have that $A \cap B = \{5\}$, so that $P(A \cap B) = 1/6$. We also have that $P(B) = 1/2$. Hence we deduce that:

$$P(A | B) = \frac{1/6}{1/2} = 1/3. \quad (2.9)$$

Note that this is indeed lower than the probability of A , which was initially $P(A) = 1/2$.

2.2.6 Conditional probability and independence

As conditional probabilities allow us to reason about how probability change when we acquire knowledge about another event, let’s try to use them and quantify our intuition about independence. Our intuition about two events A and B being independent

corresponded to the fact that knowing about one event did not affect the other. In particular, the probability of A should be the same as that of A given B . We will prove this fact.

From the definition of conditional probability (eq. (2.8)), we have that:

$$P(A | B) = \frac{P(A \cap B)}{P(B)}. \quad (2.10)$$

Now, we have in addition by the definition of independence (eq. (2.7) that $P(A \cap B) = P(A)P(B)$. Substituting this in the equation above, we have that:

$$P(A | B) = \frac{P(A \cap B)}{P(B)} = \frac{P(A)P(B)}{P(B)} = P(A), \quad (2.11)$$

where we have cancelled $P(B)$ in the numerator and denominator.

2.2.7 Conditional probability and additivity

Conditional probabilities verify most of the properties of regular probabilities, and in particular the law of additivity for disjoint events and the inclusion exclusion rule. We prove the law of additivity of disjoint event for conditional probabilities.

Let A and B be disjoint events, that is, $A \cap B = \emptyset$. Let C be any event with $P(C) > 0$. We would like to show that

$$P(A \cup B | C) = P(A | C) + P(B | C). \quad (2.12)$$

We substitute the definition of conditional probabilities, and have to show that

$$\frac{P((A \cup B) \cap C)}{P(C)} = \frac{P(A \cap C)}{P(C)} + \frac{P(B \cap C)}{P(C)}. \quad (2.13)$$

Cancel the denominators to obtain

$$P((A \cup B) \cap C) = P(A \cap C) + P(B \cap C). \quad (2.14)$$

Now, we would like to simplify the expression $(A \cup B) \cap C$. The following formulae, known as distributivity property, gives us the correct expressions:

Lemma 1 (Distributivity) *For any events A, B, C , we have that*

$$(A \cup B) \cap C = (A \cap C) \cup (B \cap C), \quad (2.15)$$

$$(A \cap B) \cup C = (A \cup C) \cap (B \cup C). \quad (2.16)$$

Hence we may simplify $(A \cup B) \cap C = (A \cap C) \cup (B \cap C)$. Now, we would like to apply the additivity rule. However, we first need to ensure that those two sets are disjoint, i.e. have empty intersection. We can check that

$$(A \cap C) \cap (B \cap C) = A \cap C \cap B \cap C = (A \cap B) \cap C = \emptyset \cap C = \emptyset, \quad (2.17)$$

hence we may apply the additivity rule, which gives us

$$P((A \cap C) \cup (B \cap C)) = P(A \cap C) + P(B \cap C). \quad (2.18)$$

2.2.8 Law of total probability

In some cases, we may want to compute the probability of an event where we only know conditional probabilities of that event. For example, suppose that we wish to compute $P(A)$ knowing $P(A | B)$ and $P(A | B^C)$. The *law of total probability* gives a formula that allows us to do so.

$$P(A) = P(A | B) P(B) + P(A | B^C) P(B^C) \quad (2.19)$$

Intuitively, we know that either B or its complement B^C happen. So understanding what happens to A under both cases should allow us to reconstruct what happens to A , by weighting by the probability that each scenario happens.

Let us prove the formula. Note that we have by definition of the conditional probability that $P(A | B) P(B) = P(A \cap B)$, and similarly $P(A | B^C) P(B^C) = P(A \cap B^C)$. Now, we have again that $A \cap B$ and $A \cap B^C$ are disjoint, hence we may apply the additivity rule to obtain that the right hand side is equal to $P((A \cap B) \cup (A \cap B^C))$. Now applying the distributivity rule in reverse, this is equal to $P(A \cap (B \cup B^C)) = P(A \cap \Omega) = P(A)$, as required.

2.2.9 Conditional probabilities and bayes rule

Given two events A and B , we will often be interested in relating $P(A | B)$ and $P(B | A)$. For example, suppose that we are testing a medical diagnosis to screen for a disease, call it disease Z . We are given the following facts about the test: for a person that has disease Z , the test will return positive 99% of the time (in other words, we have a false negative rate of 1%). For a person that does not have disease Z , the test will return positive 1% of the time (in other words, we have a false positive rate of 1%).

Suppose that we administer the test to a patient that just came to the hospital, and the result is positive. What is the probability that the patient has the disease? A tempting answer is simply that the probability is 99%, after all, the test is only wrong 1% of the time. However, we will see that this answer is in fact quite wrong.

Let's set this problem up in a mathematical fashion, and let T be the event that the test is positive (so that T^C is the event that the test is negative), and let D be the event that the person has disease Z (and hence D^C the event that the person is healthy). Putting the given assertions about the test formally, we have that

$$\begin{aligned} P(T | D) &= 0.99, \\ P(T | D^C) &= 0.01. \end{aligned}$$

We are looking to answer the question: what is the probability that someone who tested positive has disease Z . Formally, we would like to compute $P(D | T)$.

Baye's rule gives a formula to compute this quantity, by

$$P(D | T) = \frac{P(T | D) P(D)}{P(T)}. \quad (2.20)$$

Let's apply the formula to the example first. We are given $P(T | D)$ from the problem statement, but we are neither given $P(D)$ nor $P(T)$. Hence, suppose that $P(D) = 0.01$, or that one in a hundred person has this disease Z . We may now compute $P(T)$ according to the law of total probability (eq. (2.19)) to obtain

$$\begin{aligned} P(T) &= P(T | D) P(D) + P(T | D^C) P(D^C) \\ &= 0.99 \times 0.01 + 0.01 \times 0.99 \\ &= 0.198 \end{aligned}$$

Plugging this quantity into the Bayes rule (eq. (2.20)), we obtain that

$$P(D | T) = \frac{0.99 \times 0.01}{0.198} = 0.5 \quad (2.21)$$

Hence a person testing positive only has about a 50-50 chance of actually having disease Z !

Let's finish by proving the Bayes rule. By definition of the conditional distribution, we have that $P(T | D) = P(T \cap D) / P(D)$. Hence we have that the right hand side is:

$$\frac{P(T \cap D) P(D)}{P(D) P(T)} = \frac{P(T \cap D)}{P(T)} = P(D | T). \quad (2.22)$$

2.3 Random variables

Random variables are the main probability tools we will be using to discuss data. They enable us to apply the notions of probability we learned on numeric data often encountered in statistics. Random variables represent numerical outcomes that are random or unknown. We will be interested in how to characterise the randomness of these numerical quantities.

In statistics, we will often use randomness to model processes that may not necessarily be inherently random, but are too complex to understand fully. For example, consider the problem of forecasting weather. One might argue that if we had an extremely powerful computer, and as many sensors as we desire, it would be possible to exactly predict the weather. However, within our knowledge and technology today, this is not possible. Hence we model the outcome as random. This is an example of *epistemic* randomness. Another example of such randomness is the coin toss: one could argue that with extremely precise sensors, and using Newton's law of motions, one could predict exactly how the coin lands. On the other hand, we believe that some processes are inherently random, which we refer to as *ontic* randomness. For example, the radioactive decay of uranium is understood to be inherently random according to the current theory of quantum mechanics.

We will be using the formalism of probability and random variables to model both the inherent ontic randomness, and (most often) epistemic randomness, whether due to incomplete information or complexity of the process. The calculus of probability is the same whether the randomness is ontic or epistemic.

2.3.1 What is a random variable?

A random variable is a numerical variable whose value is random. Random variables can be used to model the outcome of experiments or other phenomenon that have some randomness, and will usually be denoted X, Y, Z , etc. For example, the value of the FTSE 100 Index tomorrow can be viewed as a random variable, as it is unknown today and has some randomness. Similarly, the temperature tomorrow could also be modelled as a random variable. Random variables will often also be appropriate to model information that is known, but merely hidden from us. For example, one could model a patient's blood pressure as a random variable. Although that information is not random – indeed we could measure it – it is unknown to us, and falls under epistemic randomness.

2.3.2 Random variables and events

We would like to tie the notion of random variables back to the theory of probability we have developed in the previous section, which was centered around the notion of events. How do we obtain events from random variables? As a random variable is a random quantity, any assertion concerning that variable is an event.

Example 6 Weather forecasting Let X be a random variable representing the temperature tomorrow in degrees Celsius. Then, $X < 0$ is an event, and corresponds to the colloquial event that the temperature will be below freezing. In particular, we may form the question “what is the probability that the temperature will be below freezing tomorrow?”. That probability would be written $P(X < 0)$.

In principle, knowing the probability of all the events concerning a random variable would completely characterise the randomness of that quantity. However, there are many (usually infinitely many) such events, and describing them can be difficult. Instead, we will see that we can usually characterise the randomness completely with much less information.

2.3.3 Discrete random variables

We say a random variable is discrete if it takes finitely many (or at most countably many values). Most often, a discrete random variable will take integer values. For example, the number of heads in 10 coin tosses is a discrete random variable taking integer values between 0 and 10. Another example is the number of customers that visited a store in a given day.

Discrete random variables can be characterised completely by the so-called probability mass function, which describes the probability of each possible value of a discrete random variable. We define the p.m.f. of X to be the function f_X , where

$$f_X(x) = P(X = x). \quad (2.23)$$

For example, if X is the number of heads in 10 coin tosses, then $f_X(1)$ is the probability that we obtain *exactly* one head in 10 tosses.

Example 7 Two coin tosses Let us compute the p.m.f. for X , where X is the number of heads in two coin tosses. For the case of two coins, there are four possible outcomes, being HH , HT , TH , TT , each equally likely. Hence we have that

$$\begin{aligned} f_X(2) &= P(X = 2) &&= \frac{1}{4}, \\ f_X(1) &= P(X = 1) &&= \frac{1}{2}, \\ f_X(0) &= P(X = 0) &&= \frac{1}{4}. \end{aligned}$$

The p.m.f. of a random variable allows us to compute the probability of all events concerning that random variable. For example, for the case of two tosses, consider the event $X \geq 1$ (at least one of the toss lands on heads). Then, note that $X \geq 1$ is equivalent to saying $X = 1$ or $X = 2$, and these two events are disjoint. We may thus apply the additivity rule to obtain

$$P(X \geq 1) = P(X = 1) + P(X = 2) = f_X(1) + f_X(2) = \frac{3}{4}. \quad (2.24)$$

In this sense, to define a discrete random variable, it suffices to give its p.m.f.

As the p.m.f. encodes probabilities, it inherits some properties of probabilities. In particular, we have that $0 \leq f_X(x) \leq 1$ for all possible values of x . Furthermore, we have that:

$$\sum_x f_X(x) = 1, \quad (2.25)$$

where the sum is taken over all possible values of X . This encodes the fact that X has to take some value, and corresponds to the fact that $P(\Omega) = 1$. We say that the p.m.f. is *normalized*.

2.3.4 Expectation

In addition to understanding the likelihood of certain outcomes, we will often be interested in capturing the average value of a random variable, also called its expectation. We define the expectation of a discrete random variable X with p.m.f. f_X to be:

$$\mathbb{E} X = \sum_x x P(X = x) = \sum_x x f_X(x) \quad (2.26)$$

where the sum is over all possible values of X .

Example 8 Two coin tosses Let X be the number of heads in two coin tosses, and recall the p.m.f. of X from the previous section. We would like to compute $\mathbb{E} X$. X can take values 0, 1 or 2, hence plugging the values into the formula yields:

$$\mathbb{E} X = 0 \times \frac{1}{4} + 1 \times \frac{1}{2} + 2 \times \frac{1}{4} = 1. \quad (2.27)$$

Thus the average number of heads in two coin tosses is 1, as expected.

2 Probability

Now, in addition to expectation of random variables, we will also be interested in computing expectations of functions of random variables. We define the expectation of a function $g(X)$ as follows:

$$\mathbb{E} g(X) = \sum_x g(x) \mathbb{P}(X = x) = \sum_x g(x) f_X(x). \quad (2.28)$$

Example 9 Two coin tosses Let X be the number of heads in two coin tosses, and recall the p.m.f. of X from the previous section. Let $g(x) = x^2$, and suppose that we would like to compute $\mathbb{E} g(X)$. We write

$$\mathbb{E} g(X) = \mathbb{E} X^2 = 0^2 \times \frac{1}{4} + 1^2 \times \frac{1}{2} + 2^2 \times \frac{1}{4} = \frac{3}{2}. \quad (2.29)$$

Note that $\mathbb{E} X^2$ is not equal to $(\mathbb{E} X)^2$.

2.3.5 Population statistics

The notion of expectations allows us to define population statistics of a random variable. For example, we may define the mean, or the variance, of a random variable. To do so, we will often replace the average that appears in the statistic (i.e. the $n^{-1} \sum$) by the expectation sign.

For example, we have defined the sample mean as

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n X_i, \quad (2.30)$$

hence we replace the sum sign to obtain that the sample mean (also called expectation) should simply be

$$\mathbb{E} X. \quad (2.31)$$

Similarly, we have define the sample variance as

$$\frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2. \quad (2.32)$$

Now, we replace the sample mean that appears above, and the sum sign, to obtain that the variance of a random variable is defined as

$$\sigma^2 = \mathbb{E}(X - \mathbb{E} X)^2 \quad (2.33)$$

Note that the population mean and variance are no longer random variables, but simple numbers instead.

Example 10 Two coin tosses Let X be the number of heads in two coin tosses. We recall that we had computed $\mathbb{E} X = 1$. Let's compute the variance of X .

$$\begin{aligned}
\sigma^2 &= \mathbb{E}(X - \mathbb{E} X)^2 \\
&= \mathbb{E}(X - 1)^2 \\
&= (0 - 1)^2 \times \frac{1}{4} + (1 - 1)^2 \times \frac{1}{2} + (2 - 1)^2 \times \frac{1}{4} \\
&= \frac{1}{2}.
\end{aligned}$$

Hence the variance of the number of heads in two coin tosses is $\sigma^2 = 1/2$.

2.3.6 Continuous random variables

As opposed to discrete random variables, *continuous* random variables take values over an interval of the real line. For example, a continuous random variable is an appropriate model for the number of seconds (and fractions of a second) between two radioactive decay events of a sample of uranium. Another example of a continuous random variable might be the temperature tomorrow.

Although continuous random variables behave similarly to discrete random variables, and we will be using similar tools to understand both, they can be mathematically more delicate due to the following fact: a continuous random variable is equal to a given number with probability 0! Indeed, if X is a continuous random variable, then we must have that $P(X = x) = 0$ for all x real. In particular, we won't be able to define a p.m.f. for a continuous random variable, as it would simply be 0 everywhere. However, we can still define a very similar object by making use of calculus.

For continuous random variables, the appropriate object we wish to define is the so-called *density* probability density function, also called *p.d.f.*, which encodes the *infinitesimal* probability of the random variable being equal to a given value. For example, let X be the random variable that corresponds to a (uniformly) random number on the interval $[0, 1]$ (i.e. the set of all the real numbers from 0 to 1). We say X follows the uniform distribution on $[0, 1]$. Now, although we have that $P(X = 0.5) = 0$, as it essentially never happens that the random number is *exactly* 0.5, we can certainly compute

$$P(0 \leq X \leq 0.5) = \frac{1}{2}. \quad (2.34)$$

Indeed, as X is uniformly random, it certainly has a 50-50 chance of being in either the first or the second half of the interval.

Now, we can extend this idea to compute for example $P(0 \leq X \leq 0.25) = 0.25$, as X has 1/4 chance of being in the first quarter. In general, we have for any $0 < h < 1$ that:

$$P(0 \leq X \leq h) = h. \quad (2.35)$$

Hence we could consider taking a limit

$$\lim_{h \rightarrow 0} \frac{1}{h} P(0 \leq X \leq h) = 1, \quad (2.36)$$

which corresponds to the idea of a density function.

2.3.7 Distribution and density functions

In order to make the previous idea more rigorous and precise, we will take a detour to define the (cumulative) *distribution* function. If X is a random variable (continuous or discrete), we may define the distribution function F_X :

$$F_X(x) = P(X \leq x). \quad (2.37)$$

For example, we have computed previously that for a uniform random variable on $[0, 1]$, we have

$$F_X(x) = x, \text{ for } 0 \leq x \leq 1. \quad (2.38)$$

The distribution function similarly fully characterises the distribution of a random variable, and allows us to compute probability of events involving the random variable.

Example 11 Let X be as above with the uniform distribution on $[0, 1]$. Suppose we wish to compute $P(1/4 \leq X \leq 3/4)$. Note that we can see that this is $1/2$, as this interval occupies half of the space. Let us also obtain the result using the distribution function.

Now, note that $X \leq 3/4$ is the same as $X \leq 1/4$, or $1/4 \leq X \leq 3/4$, and the latter two are disjoint. Hence we may apply the additivity rule to obtain

$$P(X \leq 3/4) = P(X \leq 1/4) + P(1/4 \leq X \leq 3/4), \quad (2.39)$$

which we may re-arrange to obtain:

$$P(1/4 \leq X \leq 3/4) = P(X \leq 3/4) - P(X \leq 1/4) = F_X(3/4) - F_X(1/4) = 1/2. \quad (2.40)$$

Although the distribution function can perfectly encode the information of the distribution of a random variable, it is sometimes inconvenient, for example, it cannot be directly used to compute expectations.

Let us go back to the problem of defining a density for a continuous random variable. As we have remarked before, asking for X to be exactly equal to some value x always has probability 0. However, suppose instead that we just wished X to be very close, say within $h/2$, where $h > 0$. Then, we would have that

$$P(x - h/2 \leq X \leq x + h/2) = F_X(x + h/2) - F_X(x - h/2). \quad (2.41)$$

As before, suppose that we normalize by h , and take the limit as $h \rightarrow 0$. We would then have that

$$\lim_{h \rightarrow 0} \frac{1}{h} P(x - h/2 \leq X \leq x + h/2) = \lim_{h \rightarrow 0} \frac{1}{h} (F_X(x + h/2) - F_X(x - h/2)) = F'_X(x), \quad (2.42)$$

where we have recognised the expression for the derivative of F_X at x on the right hand side. We thus define the (probability) density function f_X of X to be:

$$f_X(x) = F'_X(x). \quad (2.43)$$

Example 12 As above, let X be the uniform random variable on $[0, 1]$. We had computed previously that $F_X(x) = x$, for $0 \leq x \leq 1$. Hence we can compute the density function to be:

$$f_X(x) = \frac{d}{dx}x = 1, \text{ for } 0 \leq x \leq 1. \quad (2.44)$$

We will most often only be given an expression for the density f_X , and not the distribution F_X , as the former tends to be mathematically simpler. How can we use f_X to compute probability of events? Let X be a random variable, and suppose that we only knew f_X but wished to compute $P(0 \leq X \leq 1)$. From our previous discussion, we know that we have

$$P(0 \leq X \leq 1) = F_X(1) - F_X(0), \quad (2.45)$$

but we do not have access to F . However, by the fundamental theorem of calculus, we know that

$$\int_0^1 f_X(x)dx = F_X(1) - F_X(0), \quad (2.46)$$

hence we deduce that we have

$$P(0 \leq X \leq 1) = \int_0^1 f_X(x)dx \quad (2.47)$$

Example 13 As above, let X be the uniform random variable on $[0, 1]$. We had computed $f_X(x) = 1$, so we have that

$$P(1/4 \leq X \leq 3/4) = \int_{1/4}^{3/4} f_X(x)dx = \frac{1}{2}. \quad (2.48)$$

Finally, as the p.d.f. f_X is related to probabilities, it must obey similar rules. In particular, we have that $f_X(x) \geq 0$ for all x . However, it is not necessary that $f_X(x) \leq 1$, as f_X itself is not a probability. We do have that:

$$\int_{-\infty}^{\infty} f_X(x)dx = 1, \quad (2.49)$$

and say that f_X must be normalized.

2.3.8 Expectation (bis)

We may now define expectations for continuous random variable. In general, the p.d.f. we play the role of the p.m.f. for a continuous random variable, and we will have to replace summation with integration. Thus, if X is a continuous random variable with p.d.f. f_X , we may define its expectation $\mathbb{E} X$ as:

$$\mathbb{E} X = \int x f_X(x)dx, \quad (2.50)$$

where the integral is over all possible values of X (usually, an interval or the real line).

2 Probability

Example 14 Uniform random variable Let us compute the expectation of a uniform random variable X on $[0, 1]$. Recall that we have $f_X(x) = 1$ for $0 \leq x \leq 1$. Hence we have that

$$\mathbb{E} X = \int_0^1 x \cdot 1 dx = \frac{1}{2}. \quad (2.51)$$

Note that the bound of integrations correspond to the possible values of x .

We may similarly define the expectation of functions of continuous random variables. Let $g(x)$ be a function, we may then define $\mathbb{E} g(X)$ to be

$$\mathbb{E} g(X) = \int g(x) f_X(x) dx, \quad (2.52)$$

where again the integral is over all possible value of X .

Example 15 Let us compute the expectation of the function $g(x) = x^2$ of a uniform random variable X on $[0, 1]$. Recall that we have $f_X(x) = 1$ for $0 \leq x \leq 1$. Hence we have that:

$$\mathbb{E} X^2 = \int_0^1 x^2 \cdot 1 dx = \frac{1}{3}. \quad (2.53)$$

Now that we have the expectation of a continuous random variable, this allows us to define the variance of continuous random variable using the exact same formula as eq. (2.33), namely:

$$\sigma^2 = \mathbb{E}(X - \mathbb{E} X)^2 \quad (2.54)$$

Example 16 Let us compute the variance of a uniform random variable X on $[0, 1]$. Recall that we have $f_X(x) = 1$ for $0 \leq x \leq 1$, and that $\mathbb{E} X = 1/2$. Hence we have that:

$$\sigma^2 = \int_0^1 (x - 1/2)^2 \cdot 1 dx = \frac{1}{12} \quad (2.55)$$

2.4 Common distributions

It will often be useful to have distributions for which we understand the characteristics (such as p.m.f./p.d.f., mean, variance, etc.) well to model different types of variables. In this section, we describe some common distributions, their probabilistic properties, and some of the most common uses.

Most often, we will describe a *family* of distributions, that is, a set of distributions parametrised by some real parameters. These parameters will describe the exact behaviour of the distribution.

2.4.1 Bernoulli distribution

The *Bernoulli* distribution is a discrete distribution that represents a coin flip as a 0 – 1 outcome. It is parametrised by p , the probability that the outcome is 1, where $0 \leq p \leq 1$. If X has the Bernoulli distribution with parameter p , we will write $X \sim \text{Bernoulli}(p)$.

The p.m.f. of the bernoulli distribution is given by:

$$P(X = 0) = 1 - p \text{ and } P(X = 1) = p \quad (2.56)$$

We have that it has expectation $\mathbb{E} X = p$ and variance $\sigma^2 = p(1 - p)$.

The Bernoulli distribution is often used to model binary outcomes, for example, whether a patient was cured by a drug, or whether the customer purchased the product.

2.4.2 Binomial distribution

The *Binomial* distribution is a discrete distribution that represents a series of coin flip by counting the number of heads. It is parametrised by p , the probability that the flip is a head, and by n , the number of flips. We write $X \sim \text{Binom}(n, p)$.

The p.m.f. of the Binomial distribution is given by:

$$P(X = k) = \binom{n}{k} p^k (1 - p)^{n-k} \quad (2.57)$$

We have that it has expectation $\mathbb{E} X = np$ and variance $\sigma^2 = p(1 - p)$.

A binomial distribution is the sum of n independent Bernoulli distribution. It is most often used to model sums of binary outcomes. For example, the number of patient for which the drug trial was a success, or the number of match wins by a Basketball team in a season.

Note that if $X \sim \text{Binom}(n_1, p)$, and $Y \sim \text{Binom}(n_2, p)$, then $X + Y \sim \text{Binom}(n_1 + n_2, p)$.

2.4.3 Poisson distribution

The *Poisson* distribution is a discrete distribution that is usually used to represent a count. It is parametrised by a rate parameter λ , where $\lambda > 0$. We write $X \sim \text{Poisson}(\lambda)$.

The p.m.f. of the Poisson distribution is given by:

$$P(X = k) = \frac{1}{k!} \lambda^k e^{-\lambda}. \quad (2.58)$$

The Poisson distribution has mean $\mathbb{E} X = \lambda$ and variance $\sigma^2 = \lambda$.

The Poisson distribution is commonly used to model counts of occurrences of events that are fairly independent. For example, it can be a good model for the number of customers in a store on a given day.

Note that if $X \sim \text{Poisson}(\lambda_1)$ and $Y \sim \text{Poisson}(\lambda_2)$, then $X + Y \sim \text{Poisson}(\lambda_1 + \lambda_2)$.

2.4.4 Uniform distribution

The *uniform* distribution is a continuous distribution that represents values that are “equally likely” on an interval $[a, b]$ with $a < b$. We write $X \sim \mathcal{U}[a, b]$.

The p.d.f. of the uniform distribution is given by:

$$f_X(x) = \frac{1}{b - a}, \text{ for } a \leq x \leq b, \quad (2.59)$$

2 Probability

and $f_X(x) = 0$ everywhere else. The uniform distribution has mean $(a + b)/2$, and variance $(b - a)^2/12$.

The uniform distribution is often used as a simple probabilistic model for values that must be in a given interval. However, it is not particularly suitable for statistical models, as the assumption that every possible value is equally likely is usually unrealistic.

2.4.5 Normal distribution

The *Normal* distribution is a continuous distribution that is a standard model to represent a continuous quantity. It is parametrised by μ , the mean of the normal, and σ^2 , the variance of the normal.

The p.d.f. of the normal distribution is given by:

$$f_X(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left\{-\frac{1}{2\sigma^2}(x - \mu)^2\right\}. \quad (2.60)$$

The normal distribution has mean $\mathbb{E} X = \mu$ and variance σ^2 . The specific distribution with mean $\mu = 0$ and variance $\sigma^2 = 1$ is called the *standard normal* distribution. Its p.d.f. is usually written $\phi(x)$.

The normal distribution is commonly used to model any continuous quantity. For example, it can be a good model for the weight of an animal, or the temperature of a sample.

Note that if $X \sim \mathcal{N}(\mu_1, \sigma_1^2)$ and $Y \sim \mathcal{N}(\mu_2, \sigma_2^2)$, then we have that $X + Y \sim \mathcal{N}(\mu_1 + \mu_2, \sigma_1^2 + \sigma_2^2)$.

2.4.6 Exponential distribution

The *exponential* distribution is a continuous distribution on the positive real numbers. It is parametrised by the rate $\lambda > 0$. We write $X \sim \text{Exp}(\lambda)$.

The p.d.f. of the exponential distribution is given by:

$$f_X(x) = \lambda e^{-\lambda x}. \quad (2.61)$$

It has mean $\mathbb{E} X = 1/\lambda$ and variance $\sigma^2 = 1/\lambda^2$.

The exponential distribution is commonly used to model waiting times or survival times. For example, it might be a good distribution to model the time between two customers, or the time before a machine needs repairs. It is closely linked to the Poisson distribution: if e.g. customers arrive independently with a time gap distributed according to an exponential distribution, the number of customers in a unit of time is then Poisson.

2.5 Jointly distributed variables

We will often want to understand the behaviour of several related random variables at the same time. For example, we may be interested in understanding both the height and the weight of an individual. To do so, we may adapt the tools we have developed for a single random variable to the case of more than one random variable.

Die	1	2	3	4	5	6
X	0	1	0	1	0	1
Y	0	1	1	0	1	0

Table 2.1: List of outcomes for a single die roll

2.5.1 Joint distribution

We can express the random behaviour of two random variables using a *joint distribution*. For example, suppose X and Y are discrete random variables, the joint p.m.f. is then defined by

$$f_{XY}(x, y) = P(X = x, Y = y). \quad (2.62)$$

If X and Y are continuous random variables, we may instead define their joint density f_{XY} .

Example 17 Suppose that we roll a single die, and let X be the random variable such that $X = 1$ if the roll is even, and $X = 0$ otherwise. Let Y be the random variable such that $Y = 1$ if the roll is prime, and $Y = 0$ otherwise. Let us collect all the possible outcomes of the die, and the corresponding values of X and Y into table 2.1.

By examining all the possible outcomes, we may compute the following joint p.m.f. for X and Y :

$$\begin{aligned} f_{XY}(0, 0) &= P(X = 0, Y = 0) = 1/6 \\ f_{XY}(1, 0) &= P(X = 1, Y = 0) = 2/6 \\ f_{XY}(0, 1) &= P(X = 0, Y = 1) = 2/6 \\ f_{XY}(1, 1) &= P(X = 1, Y = 1) = 1/6 \end{aligned}$$

2.5.2 Joint distribution and events

Given two random variables X and Y , and their joint p.m.f. or p.d.f. f_{XY} , we may compute the probability of any events. For example, for X and Y given, we may be interested in computing $P(X \leq 1, Y \geq 0)$. However, we may also desire to compute events of the type $P(X \geq Y/2)$.

For X and Y discrete, we can simply sum over all the combinations of X and Y that verify the condition:

$$P(X \geq Y/2) = \sum_{x, y: x \geq y/2} f_{XY}(x, y). \quad (2.63)$$

For X and Y continuous, we require double integration over the set of all x and y satisfying the condition, e.g.:

$$P(X \geq Y/2) = \iint_{x, y: x \geq y/2} f_{XY}(x, y) dx dy \quad (2.64)$$

2 Probability

Example 18 Suppose X and Y are two continuous variables on $[0, 1]$, with joint p.d.f. given by $f_{XY}(x, y) = 4xy$ for $0 \leq x \leq 1$ and $0 \leq y \leq 1$. Let us compute $P(X \leq Y/2)$.

We compute the double integral:

$$\begin{aligned}
 P(X \leq Y/2) &= \iint_{x, y: x \leq y/2} 4xy \, dx \, dy \\
 &= 4 \int_0^1 \left(\int_0^{y/2} xy \, dx \right) dy \\
 &= 4 \int_0^1 y \left(\int_0^{y/2} x \, dx \right) dy \\
 &= 4 \int_0^1 yy^2/8 \, dy \\
 &= \frac{1}{2} \int_0^1 y^3 \, dy \\
 &= \frac{1}{8}
 \end{aligned}$$

2.5.3 Joint distribution and expectations

Similarly, given two random variables X and Y , and their joint p.m.f. or p.d.f. f_{XY} , we may compute the expectation of a function $g(X, Y)$ of both variables.

For discrete random variables, we sum over all possible pairs of values of x and y :

$$\mathbb{E} g(X, Y) = \sum_{x, y} g(x, y) f_{XY}(x, y). \quad (2.65)$$

For continuous random variables, we integrate over all possible values of x and y :

$$\mathbb{E} g(X, Y) = \iint g(x, y) f_{XY}(x, y) \, dx \, dy \quad (2.66)$$

Example 19 Consider X, Y random variables on $[0, 1]$ with joint p.d.f. $f_{XY}(x, y) = 6(x - y)^2$. Let us compute the expectation of $g(X, Y) = XY$.

$$\begin{aligned}
 \mathbb{E} XY &= \iint xy 6(x - y)^2 \, dx \, dy \\
 &= 6 \int_0^1 \int_0^1 xy(x - y)^2 \, dx \, dy \\
 &= 6 \int_0^1 \int_0^1 x^3y - 2x^2y^2 + xy^3 \, dx \, dy \\
 &= 6 \int_0^1 \left[\frac{1}{4}y - \frac{2}{3}y^2 + \frac{1}{2}y^3 \right] dy \\
 &= 6 \times \left[\frac{1}{8} - \frac{2}{9} + \frac{1}{8} \right] \\
 &= 12 \left(\frac{1}{8} - \frac{1}{9} \right) \\
 &= 1/6.
 \end{aligned}$$

2.5.4 Marginal distribution

Given two variables X and Y , with a joint p.m.f. or joint p.d.f. f_{XY} , suppose that we were only interested in the first variable X . Indeed, suppose that we “forget” about Y , what is the distribution of X ?

If X, Y are discrete, we may compute the p.m.f. f_X of X by

$$f_X(x) = \sum_y f_{XY}(x, y), \quad (2.67)$$

where the sum is taken over all possible values of y .

Similarly, if X, Y are continuous, we may compute the p.d.f. f_X of X by:

$$f_X(x) = \int f_{XY}(x, y) dy, \quad (2.68)$$

where the integral is taken over all possible values of y .

Example 20 As in the previous section, consider a die roll, and X the random variable which is 1 if the roll is even, and 0 otherwise, and Y , the random variable which is 1 if the roll is prime, and 0 otherwise.

Let us compute the marginal p.m.f. of X . We have

$$P(X = 0) = P(X = 0, Y = 0) + P(X = 0, Y = 1) = \frac{1}{6} + \frac{2}{6} = \frac{1}{2}. \quad (2.69)$$

2.5.5 Independent random variables

Let X and Y be two discrete random variables, and suppose that the events $X = x$ and $Y = y$ are independent for all x and y . Then, we would have by definition that:

$$P(X = x, Y = y) = P(X = x)P(Y = y). \quad (2.70)$$

That is to say, we would have in terms of the p.m.f. that

$$f_{XY}(x, y) = f_X(x)f_Y(y). \quad (2.71)$$

When eq. (2.71) holds, for the p.m.f. if the variables are discrete, or the p.d.f. when the variables are continuous, we say that the two variables are *independent*.

Example 21 In the previous example of the die roll, we compute the joint p.m.f. and the marginal p.m.f. of X . It is not hard to see that the p.m.f. of Y is given by $P(Y = 0) = P(Y = 1) = 1/2$, as the event that a roll is odd happens with probability 0.5.

Now, note that $P(X = 0, Y = 0) = 1/6$, whereas $P(X = 0)P(Y = 0) = 0.5 \times 0.5 = 1/4$, hence X and Y are *not* independent.

2.5.6 Conditional distributions

Let X and Y be two discrete random variables. We can consider the conditional distribution of X given $Y = y$, which is given by the definition of a conditional probability as:

$$P(X = x \mid Y = y) = \frac{P(X = x, Y = y)}{P(Y = y)}. \quad (2.72)$$

Hence in terms of p.m.f., the conditional p.m.f. of X given Y , written $f_{X|Y}(x \mid y)$, is given by:

$$f_{X|Y}(x \mid y) = \frac{f_{XY}(x, y)}{f_Y(y)} \quad (2.73)$$

For continuous random variables, the event $P(Y = y)$ always has probability 0, hence we may not directly define conditional distributions. However, we will define the conditional density according to eq. (2.73), replacing the p.m.f. by p.d.f.

2.6 Operation on random variables

As random variables represent numerical quantities, it is natural to operate on them as we would on numbers. The result of these operations define new random variables, for which we would like to understand their properties.

2.6.1 Transforming a random variable

Suppose that we would like to model the wealth of an individual as a random variable. This is a continuous quantity, hence a normal variable would be appropriate. However, empirical data and economic theories indicate that the more appropriate quantity to model is not the wealth, but the logarithm of the wealth. Suppose that we model the logarithm of the wealth X as a normal random variable. What is the distribution of the wealth $Y = e^X$?

In general, suppose that we have some continuous random variable X with density f_X , and some function g . We would like to compute the density of the variable $Y = g(X)$.

We have that

$$f_Y(y) = \frac{f_X(x)}{|g'(x)|}, \quad (2.74)$$

where $y = g(x)$.

Example 22 Suppose that X is standard normal, and $Y = e^X$. Then, the density of Y is given by:

$$f_Y(y) = \frac{1}{\sqrt{2\pi}} e^{-x^2/2} \frac{1}{e^x}. \quad (2.75)$$

Now, if $y = e^x$, we have $x = \log y$, hence replacing in the equation above gives:

$$f_Y(y) = \frac{1}{\sqrt{2\pi}} \frac{1}{\log y} e^{-(\log y)^2/2}. \quad (2.76)$$

This is known as the *lognormal* distribution.

2.6.2 Sums of random variables

Given two random variables X and Y , we will often be interested in understanding the behaviour of the sum $X + Y$. For example, as we have seen previously, the binomial distribution can be understood as a sum of Bernoulli random variables.

Suppose X and Y are independent discrete random variables. Then $X + Y$ is also discrete random variable, and we may characterise its behaviour by its p.m.f. We thus wish to compute $P(X + Y = k)$ for all k integers, say. Now, let us decompose the event $X + Y = k$ depending on the value of X . Suppose $X = l$, then we must have $Y = k - l$. Conversely, if $X + Y = k$, then we must have $X = l$ and $Y = k - l$ for some l . Hence the event $X + Y = k$ is a disjoint union of the events $X = l, Y = k - l$. We thus deduce by the additivity rule that:

$$P(X + Y = k) = \sum_l P(X = l, Y = k - l), \quad (2.77)$$

where the sum is over all possible values of l . Writing this in terms of the joint p.m.f., we have that

$$f_{X+Y}(k) = \sum_l f_{X,Y}(l, k - l). \quad (2.78)$$

In the case that X and Y are independent, the joint p.m.f. factorizes, and we obtain that

$$f_{X+Y}(k) = \sum_l f_X(l) f_Y(k - l). \quad (2.79)$$

The right hand side is said to be a *convolution*.

If X and Y are continuous, we may compute the p.d.f. of the sum $X + Y$ in a similar fashion as in eq. (2.78). Indeed, we have that the p.d.f. is given by

$$f_{X+Y}(z) = \int_{-\infty}^{\infty} f_{X,Y}(t, z - t) dt. \quad (2.80)$$

If X and Y are independent, we may simplify the above expression to:

$$f_{X+Y}(z) = \int_{-\infty}^{\infty} f_X(t) f_Y(z - t) dt. \quad (2.81)$$

Example 23 Let $X \sim \text{Poisson}(\lambda_1)$ and $Y \sim \text{Poisson}(\lambda_2)$ be independent Poisson random variables. What is the distribution of $X + Y$?

By eq. (2.79), we may write down the p.m.f. of the sum as

$$P(X + Y = k) = \sum_{l=0}^k P(X = l) P(Y = k - l). \quad (2.82)$$

2 Probability

Use the formula for the Poisson p.m.f. to obtain:

$$\begin{aligned}
 P(X + Y = k) &= \sum_{l=0}^k \frac{\lambda_1^l e^{-\lambda_1}}{l!} \frac{\lambda_2^{k-l} e^{-\lambda_2}}{(k-l)!} \\
 &= e^{-(\lambda_1 + \lambda_2)} \sum_{l=0}^k \frac{1}{l!(k-l)!} \lambda_1^l \lambda_2^{k-l} \\
 &= e^{-(\lambda_1 + \lambda_2)} \frac{1}{k!} \sum_{l=0}^k \binom{k}{l} \lambda_1^l \lambda_2^{k-l} \\
 &= e^{-(\lambda_1 + \lambda_2)} \frac{(\lambda_1 + \lambda_2)^k}{k!},
 \end{aligned}$$

where we have used the binomial expansion of $(\lambda_1 + \lambda_2)^k$. We now recognise that the computed p.m.f. corresponds to that of a $\text{Poisson}(\lambda_1 + \lambda_2)$ random variable, hence we have that $X + Y \sim \text{Poisson}(\lambda_1 + \lambda_2)$.

2.7 Properties of the expectation

The expectation of a random variable will be the main quantity we will be discussing in order to characterise the average behaviour of our estimators in statistics. It will thus be useful to understand some of its algebraic properties.

Let us start by mentioning a property that you will probably find obvious: the expectation of a constant. If a is a real constant (i.e. not random), then we have that $\mathbb{E} a = a$.

2.7.1 Expectation of a sum

From the definition of the expectation, we see that the expectation is essentially a sum (or integral). In particular, it is *linear*. That is, for X and Y random variables, we have that:

$$\mathbb{E}[X + Y] = \mathbb{E} X + \mathbb{E} Y. \quad (2.83)$$

In addition, if a is a real constant, we have that

$$\mathbb{E} aX = a \mathbb{E} X. \quad (2.84)$$

Indeed, recall the definition of the expectation as a sum (for discrete random variables), we have that:

$$\begin{aligned}
 \mathbb{E}[X + Y] &= \sum_{x,y} (x + y) P(X = x, Y = y) \\
 &= \sum_{x,y} x P(X = x, Y = y) + \sum_{x,y} y P(X = x, Y = y) \\
 &= \sum_x x P(X = x) + \sum_y y P(Y = y) \\
 &= \mathbb{E} X + \mathbb{E} Y.
 \end{aligned}$$

Unlike many other properties we will cover, this property does not necessitate that X and Y be independent.

Example 24 Let us derive another formula for the variance using the linearity property. We have defined variance as:

$$\sigma^2 = \mathbb{E}(X - \mathbb{E} X)^2 \quad (2.85)$$

We can expand the square to obtain the following:

$$\begin{aligned} \sigma^2 &= \mathbb{E}[X^2 - 2X\mathbb{E} X + (\mathbb{E} X)^2] \\ &= \mathbb{E} X^2 - 2\mathbb{E}[X\mathbb{E} X] + \mathbb{E}[(\mathbb{E} X)^2] \\ &= \mathbb{E} X^2 - 2\mathbb{E} X \mathbb{E} X + (\mathbb{E} X)^2 \\ &= \mathbb{E} X^2 - (\mathbb{E} X)^2. \end{aligned}$$

We have used the fact that $\mathbb{E} X$ is a constant, and hence has expectation $\mathbb{E} X$, and the linearity of the expectation.

2.7.2 Expectation of product

We now turn to the expectation of a product of random variables. Unlike the sum, this will require that the variables be independent. Let X and Y be independent random variables, then we have that:

$$\mathbb{E}[XY] = (\mathbb{E} X)(\mathbb{E} Y) \quad (2.86)$$

Indeed, if X and Y are independent, then the p.m.f. or p.d.f. factorizes. We then have:

$$\begin{aligned} \mathbb{E}[XY] &= \sum_{x,y} xy f_{X,Y}(x,y) \\ &= \sum_{x,y} xy f_X(x) f_Y(y) \\ &= \sum_{x,y} x f_X(x) y f_Y(y) \\ &= \left(\sum_x x f_X(x) \right) \left(\sum_y y f_Y(y) \right) \\ &= \mathbb{E} X \mathbb{E} Y. \end{aligned}$$

Note that we cannot bypass the requirement for independence. For example, suppose that $X = \pm 1$ with probability half for each. Then $\mathbb{E} X = 0$. However, $\mathbb{E} X X = \mathbb{E} X^2 = 1$ whereas $(\mathbb{E} X)^2 = 0$.

2.7.3 Variance of a sum

Using the previous two examples allows us to compute the variance of a sum of random variables. Let X and Y be random variables, and let σ_{X+Y}^2 be the variance. We have

2 Probability

that:

$$\begin{aligned}\sigma_{X+Y}^2 &= \mathbb{E}(X + Y - \mathbb{E}(X + Y))^2 \\ &= \mathbb{E}(X - \mathbb{E}X + Y - \mathbb{E}Y)^2 \\ &= \mathbb{E}[(X - \mathbb{E}X)^2 + (Y - \mathbb{E}Y)^2 + 2\mathbb{E}[(X - \mathbb{E}X)(Y - \mathbb{E}Y)]] \\ &= \mathbb{E}(X - \mathbb{E}X)^2 + \mathbb{E}(Y - \mathbb{E}Y)^2 + 2\mathbb{E}[(X - \mathbb{E}X)(Y - \mathbb{E}Y)] \\ &= \sigma_X^2 + \sigma_Y^2 + 2\sigma_{XY},\end{aligned}$$

where we have put $\sigma_{XY} = \mathbb{E}[(X - \mathbb{E}X)(Y - \mathbb{E}Y)]$, the *covariance* of X and Y .

If X and Y are independent, we may compute the covariance explicitly by the product rule for expectations. Indeed, we have that:

$$\begin{aligned}\sigma_{XY} &= \mathbb{E}[(X - \mathbb{E}X)(Y - \mathbb{E}Y)] \\ &= \mathbb{E}[XY - X\mathbb{E}Y - Y\mathbb{E}X + \mathbb{E}X\mathbb{E}Y] \\ &= \mathbb{E}XY - \mathbb{E}X\mathbb{E}Y - \mathbb{E}Y\mathbb{E}X + \mathbb{E}X\mathbb{E}Y \\ &= 0,\end{aligned}$$

as $\mathbb{E}XY = \mathbb{E}X\mathbb{E}Y$ for X and Y independent. If $\sigma_{XY} = 0$, we say that X and Y are *uncorrelated*. In that case, we note that we have $\sigma_{X+Y}^2 = \sigma_X^2 + \sigma_Y^2$, that is, the variance of the sum is the sum of the variance. Note that the same relationship does *not* hold for the standard deviation.

2.8 Limit theorems

Two fundamental limit theorems of probability will allow us to justify the theory of statistics to data, by showing that at least, as the amount of data we gather goes to infinity, we can be increasingly confident about our estimates.

2.8.1 Law of large numbers

The law of large numbers describes the first-order behaviour of sums of independent and identically distributed random variables. Let X_1, \dots, X_n , be independent, identically distributed (i.i.d.) random variables. The sample mean may be written as:

$$S_n = \frac{1}{n} \sum_{i=1}^n X_i. \quad (2.87)$$

As the number of samples increases, we expect this quantity to get closer to the “true value” $\mathbb{E}X$, and to be exactly equal in the limit $n \rightarrow \infty$. This is essentially the assertion of the law of large numbers.

Theorem 1 (Law of large numbers) *Let X_1, \dots, X_n be independent, identically distributed random variables. Then, (under mild regularity assumptions), we have that (with probability 1):*

$$\lim_{n \rightarrow \infty} \frac{1}{n} \sum_{i=1}^n X_i = \mathbb{E}X \quad (2.88)$$

A couple of notes on the conditions in parentheses: there are some distributions for which the randomness is so high that we cannot define a meaningful expectation. One common example is the *cauchy* distribution. In these cases, the law of large numbers does not apply.

The limit we are taking is also a random quantity, as it depends on the random variables $X_i, i = 1, 2, \dots$. However, we are claiming that in the limit, there is in fact no randomness, and the limit always takes the value $\mathbb{E}X$. In order for this claim to make sense, we have to exclude a couple of cases, in which we get “infinitely” unlucky. For example, in a series of coin tosses, it is not *impossible* to get infinitely many tails. However, this happens with probability 0.

2.8.2 Central limit theorem

We have seen in the previous section that the sample mean of i.i.d. random variables converges to the expectation in the limit. Can we characterise the speed of convergence, and the distribution of the errors?

First, let us compute the variance of the sum. The variance is an indication of how much variation around the expectation our estimate has on average. First, we note the following property of the variance: for a random variable X and a real number a , we have that:

$$\text{Var } aX = a^2 \text{Var } X. \quad (2.89)$$

Indeed, recall that the variance has the same units as X^2 . Now, we may use the formula above with our result for the variance of a sum to compute:

$$\text{Var } \frac{1}{n} \sum_{i=1}^n X_i = \frac{1}{n^2} \sum_{i=1}^n \text{Var } X_i = \frac{1}{n^2} n \text{Var } X = \frac{\text{Var } X}{n}. \quad (2.90)$$

Hence we see that the variance decays at a rate of n^{-1} . Thus, we have that the following quantity has constant variance:

$$\sqrt{n}(\bar{x} - \mathbb{E}X). \quad (2.91)$$

The central limit theorem allows us to characterise the distribution of this quantity exactly in the limit.

Theorem 2 *Central limit theorem* Let X_i be i.i.d. random variables (with finite variance), with mean μ and variance σ^2 . Let $\bar{x}_n = n^{-1} \sum_{i=1}^n X_i$ be the sample mean of the n first samples. We have that:

$$\sqrt{n}(\bar{x}_n - \mu) \xrightarrow{d} \mathcal{N}(0, \sigma^2). \quad (2.92)$$

In the case of the central limit theorem, our quantity is converging to a random distribution. We can interpret this as saying that, if we were to repeat the experiment, and compute our error $\sqrt{n}(\bar{x} - \mu)$ again, the errors would follow a normal distribution. Note that this is a universal behaviour, and does not depend on the distribution of X . This in part underlies the prevalence of the normal distribution in statistics, as it arises naturally as the distribution of errors.

2 Probability

Example 25 Suppose that a store has an average of 200 customers per day. We wish to compute the probability that the store has more than 6150 customers over a month (30 days).

Let X_i be the number of customers on day i , and suppose that $X_i \sim \text{Poisson}(200)$. We wish to compute:

$$\begin{aligned} P\left(\sum_i X_i > 6150\right) &= P\left(\frac{1}{30} \sum_i X_i > 205\right) \\ &= P\left(\frac{1}{30} \sum_i X_i - 200 > 5\right) \\ &= P\left(\sqrt{30} \left(\frac{1}{30} \sum_i X_i - 200\right) > 5\sqrt{30}\right) \end{aligned}$$

The quantity:

$$Z = \sqrt{30} \left(\frac{1}{30} \sum_i X_i - 200\right) \tag{2.93}$$

has an approximately normal distribution with mean 0 and variance 200 by the central limit theorem.

We may thus compute:

$$P(Z > 5\sqrt{30}) = 0.02640. \tag{2.94}$$

By using properties of the Poisson distribution, we may also compute the probability exactly in terms of the Poisson distribution: $\sum_i X_i \sim \text{Poisson}(30 * 200)$, and

$$P\left(\sum_i X_i > 6150\right) = 0.02637. \tag{2.95}$$

3 Sampling

The first step in any statistical analysis is the collection of data. Although we will not always have full control over this step, it is a crucial part of the statistical analysis, and must be analyzed as such. Collecting data in a sub-optimal fashion can often produce misleading results and limit the scopes of possible conclusions.

Data collection can be broadly categorised into two types: observational studies and experiments.

3.1 Experiments

In an *experiment*, the statistician is interested in understanding the relationship between some dependent variables and independent variables. To do so, the statistician will artificially control a given variable, for example whether a patient was given a drug, and observe its average effect on the variable of interest, for example the blood pressure.

3.1.1 Randomized experiment

Most commonly, the statistician will seek to assign patients randomly to either a trial or control group. This random assignment ensures that on average, a person given the treatment is the “same” as a person not given the treatment. Mathematically, the person given a treatment has the same distribution as a person not given the treatment.

This is the most powerful type of experiment, as it ensures that the difference observed between the two groups is entirely due to the treatment assignment. There can be no *confounding* variable.

3.1.2 Control groups and placebo

Although the randomized experiment design ensures that the difference in outcome only comes from whether the patient was assigned to a control or trial group, this is usually not sufficient. Indeed, the only conclusion we would be able to rigorously obtain on that is that difference in the entire treatment (e.g. meeting with doctors, taking the drug), produced the difference in the outcome. However, we are more often interested in a more specific aspect, for example, whether the drug itself produced a difference in the outcome.

Although this difference may seem trivial, the now well-documented *placebo* effect states that the simple act of meeting with doctors, and taking a pill (which could have no medical effect), is enough to affect the outcome of a patient. In fact, the knowledge by the doctor (not the patient!) that there is a difference in treatment is enough to produce a different result. To address this issue, the clinical community has adopted the *double*

3 Sampling

blind trial, where each patient is issued either the real drug or a placebo, and the doctors do not know which.

In general, this issue highlights the fact that one should attempt to isolate the effect to be tested for as best as possible. In order to make the experiment unlikely to be affected by other possibilities, it is best to keep control and trial groups as similar as possible.

3.1.3 Sub-populations and inductive inference

A properly conducted randomized trial will usually be effective to isolate the desired effect. However, this effect is only observed for the population that participated in the study, and it is not always correct to extend it to a more general population. For example, suppose that a drug trial was conducted with 100 participants, who are white of germanic descent. However, the effect may not be the same in a hispanic population. It will often be necessary to replicate the experiment with other populations to ensure that the conclusion can be generalized.

3.2 Observational studies

It will often be the case that it is impossible or unethical to conduct the desired experiment. In this case, the statistician can only rely on an *observational* study, that is, observing different patients who happen to fall in a control or trial group. However, as the assignment to each group (for example, smoking) is no longer random, there is no guarantee that patients in one group are “statistically similar” to those in the other group. For example, there are well document links between smoking status and socio-economic status, which has further implications towards other health-related aspects, such as diet or access to healthcare.

We will thus often need to take special care in order to ensure that the comparison across groups that we wish to compute is valid, and will often not be able to reach conclusions that are quite as strong as in the experimental case. On the other hand, observational studies represent the vast majority of the data collected today, and allows us to use existing datasets instead of specifically designing an experiment for the problem at hand.

3.2.1 Prospective and retrospective studies

Observational studies can be broadly categorized into two main forms: *prospective* and *retrospective* studies. In a prospective study, the experimenter would identify individuals to observe, and collect information as time unfolds. Although these studies are able to provide strong evidence, they tend to be expensive. A famous prospective observational study is the Framingham heart study, which began in 1948 and tracked about 5000 adults and their descendants (it is today in its third generation of participants).

Retrospective studies, on the other hand, collects the data after the event of interest has taken place. For example, a researcher may identify 100 patients affected by lung cancer, and investigate their medical history. Retrospective studies are often much cheaper than

prospective studies, and are particularly adapted for studying rare events. For example, the lifetime risk of lung cancer is about 7%. A prospective study on a cohort of 100 people would thus only observe about 7 cases of lung cancer on average, which is not adequate if one wishes to study lung cancer. On the other hand, a retrospective study would be able to select post-hoc 100 people who have experienced lung cancer.

3.2.2 Confounding

Observational studies face significant problems from confounding, which is the situation that a third variable is correlated with both the dependent variable and the independent variable. For example, consider an observational study of the effectiveness of sunscreen to prevent skin cancer. If we were only to measure the usage of sunscreen and the prevalence of skin cancer, we could possibly obtain a result that those who use sunscreen tend to be more likely to obtain skin cancer. However, we did not account for exposure to the sun, the confounding variable in this case: those who are more exposed to the sun are both more likely to use sunscreen, and more likely to experience skin cancer.

3.2.3 Natural experiments

In some cases, it may be possible to interpret a mechanism outside of the experimenter's control as a randomized experiment. This idea is of particular importance in fields such as econometrics, where it is often known as an *instrumental variable*. For example, a number of current estimates in inheritability of various characteristics comes from studies of twins separated at birth.

Another famous example is that of lifetime earnings of veterans of the Vietnam war. An initial simple estimate (by comparing American adults who were deployed vs. those who were not deployed) of the effect of the war on their lifetime earnings seemed to indicate that veterans experienced a *positive* effect on their lifetime earnings. However, this comparison omits a number of effects, due to people volunteering. In 1990, Angrist used the fact that the draft lottery as a natural experiment to provide an estimate that in fact, the earnings of veterans was about 15% less than the non-veterans.

4 Estimation

The problem of estimation is one of the central problems of statistics, and is the process through which we make inferences about a population from a collected sample.

4.1 Models and likelihood

In order to talk about a population in a mathematical fashion, we require a statistical *model* of the population. This model relates parameters of interest (for example, the effectiveness of a drug) to the data we observe (the blood pressure of individuals before and after taking the drug). Mathematically, we will make use of a family of probability distribution, most commonly one that we saw in section 2.4. This family of probability distribution thus relates the parameter to the data that is observed.

Example 26 Suppose we wish to understand the bias of a coin. In this case, the parameter of interest is the bias of the coin, which may be described as a number p , where p is the probability that the coin lands on heads.

Suppose we carry out the “experiment” of flipping the coin 10 times, then a statistical model for the experiment could be to say that the number of heads observed in the experiment follows a binomial distribution, with parameters $n = 10$ and p the bias.

We saw that the p.m.f. or p.d.f. was very useful in understanding the randomness of a quantity. Conversely in estimation, it will be useful to understand the relation between the parameter and the data. However, in probability, we think of the p.m.f. or p.d.f. as a function of the outcome (the value of the random variable) for a given parameter. In estimation, we are usually given the outcome as observed in the experiment we carried out, and instead wish to understand its implications on the parameter. We thus define the *likelihood*, which is the p.m.f. or p.d.f. viewed as a function of the parameter, and is usually written $L(\theta)$, where θ is the parameter of interest.

Example 27 Continuing from the experiment of tossing 10 coins, the p.m.f. of the outcome is given by (for p fixed):

$$f_X(k) = P(X = k) = \binom{n}{k} p^k (1 - p)^{n-k} \quad (4.1)$$

We will instead view this as a function of p . Suppose that we have k fixed (for example $k = 4$ if we observed four heads in the experiment), and consider the likelihood (written $L(p)$):

$$L(p) = \binom{n}{k} p^k (1 - p)^{n-k} \quad (4.2)$$

We will usually use the *log-likelihood*, which is defined as the natural log of the likelihood, and written $\ell(\theta)$. This is mostly mathematically convenient, and will make the analysis much simpler.

4.2 Estimators

In this section, we will study *estimators*, which are functions of the data that attempt to “guess” the value of the parameter. For a given parameter θ , we will usually write $\hat{\theta}$ for an estimator of θ . Note that estimators depend on the data: they are thus *random* quantities. We will thus use the tools of probability to discuss estimators.

4.2.1 Unbiased estimators

For a general estimator, a natural question is to quantify how “good” an estimator is. A first attempt at capturing such a notion can be the very simple question: is our estimator right on average? That is, if we were to repeat the experiment many times, would the average outcome of the estimator be the true value?

We say an estimator $\hat{\theta}$ of θ is *unbiased* for θ if the following holds:

$$\mathbb{E}_{\theta} \hat{\theta} = \theta. \quad (4.3)$$

Here, the expectation is taken with respect to the data, supposing that the true parameter is given by θ .

Example 28 Consider again the case of the coin toss, and suppose that we consider the estimator:

$$\hat{p} = \frac{1}{n}X \quad (4.4)$$

Then, we claim that \hat{p} is unbiased for X . Indeed, we may compute its expectation to obtain:

$$\mathbb{E}_p \hat{p} = \mathbb{E}_p \frac{1}{n}X = \frac{1}{n} \mathbb{E}_p X = \frac{1}{n}np = p. \quad (4.5)$$

However, an unbiased estimator only measures the fact that it is as likely to overestimate as underestimate the parameter of interest. In addition, we would hope that our estimator is close to the true value on average.

If $\hat{\theta}$ is unbiased for θ , then we have that the average squared-distance to the true value is given by:

$$\mathbb{E}_{\theta}(\hat{\theta} - \theta)^2 = \mathbb{E}_{\theta}(\hat{\theta} - \mathbb{E} \hat{\theta})^2 = \text{Var } \hat{\theta}. \quad (4.6)$$

In particular, the notion of closeness for unbiased estimator may be captured by low variance of the estimator. To quantify this, we call the standard deviation of the estimator the *standard error*.

4.2.2 Mean-squared error

We will see that unbiasedness is not necessarily a desirable property of an estimator in general. Indeed, we will simply most of the time require that our estimator be close on average, or have low mean-squared error (often written *mse*):

$$\text{mse}(\theta) = \mathbb{E}_\theta(\hat{\theta} - \theta)^2. \quad (4.7)$$

The mse quantifies the average distance of our estimate from the true value. It can be decomposed into two parts for which we will give individual interpretation:

$$\begin{aligned} \text{mse}(\theta) &= \mathbb{E}_\theta(\hat{\theta} - \theta)^2 \\ &= \mathbb{E}_\theta(\hat{\theta} - \mathbb{E}_\theta \hat{\theta} + \mathbb{E}_\theta \hat{\theta} - \theta)^2 \\ &= \mathbb{E}_\theta(\hat{\theta} - \mathbb{E}_\theta \hat{\theta})^2 + (\mathbb{E}_\theta \hat{\theta} - \theta)^2 \\ &= \text{Var}_\theta \hat{\theta} + (\mathbb{E}_\theta \hat{\theta} - \theta)^2 \end{aligned}$$

We call $\text{Var}_\theta \hat{\theta}$ the variance of the estimator, and $\mathbb{E}_\theta[\hat{\theta} - \theta]$ the bias of the estimator. We thus see that we may write:

$$\text{mse} = \text{variance} + \text{bias}^2, \quad (4.8)$$

the *bias-variance* decomposition of the mse.

We will see that it can often be advantageous to trade off these two quantities, and in particular, it can sometimes be useful to incur a slight bias to reduce the variance greatly.

4.3 Maximum likelihood estimation

Maximum likelihood estimation (or *mle*) is a general method of obtaining estimators for a given quantity. For a given statistical model and likelihood, the idea is to obtain the value of the parameter that maximizes the likelihood of the data. Conceptually, our guess for the data represents is the parameter that produces the observed data with the highest probability.

We illustrate this principle through two examples: a binomial example, and an exponential example.

4.3.1 Example: binomial model

Let X follow a binomial distribution with parameters n (known) and p (parameter of interest). The likelihood is then given by:

$$L(p) = \binom{n}{k} p^k (1-p)^{n-k} \quad (4.9)$$

We wish to compute the value of p that maximizes L , which is usually referred to as the arg-max and written $\arg \max_p L(p)$. However, as \log is an increasing function, this is

4 Estimation

equivalent to computing the value of p that maximises $\log L(p) = \ell(p)$, the log-likelihood.

$$\ell(p) = k \log p + (n - k) \log(1 - p) + \log \binom{n}{k} \quad (4.10)$$

Now, note that the last term $\log \binom{n}{k}$ does not depend on p , hence can be ignored for the purpose of maximization. We are thus left to maximize $k \log p + (n - k) \log(1 - p)$. In order to do so, we will simply compute the derivative and set it to 0.

$$\frac{\partial \ell}{\partial p} = \frac{k}{p} - \frac{n - k}{1 - p} = 0. \quad (4.11)$$

Solving the above equation in p gives $\hat{p} = k/n$, the mle for p in the binomial model.

4.3.2 Example: exponential model

Suppose now instead that we have X_1, \dots, X_n i.i.d. random variables distributed with an exponential distribution of parameter λ . The likelihood is given by the joint p.d.f., which in this case is simply the product of the p.d.f. for each observation X_i as the observations are independent.

$$L(\lambda) = \prod_{i=1}^n \lambda e^{-\lambda x_i} = \lambda^n e^{-\lambda \sum_{i=1}^n x_i} \quad (4.12)$$

We may similarly compute the log-likelihood to obtain:

$$\ell(\lambda) = n \log \lambda - \lambda \sum_{i=1}^n x_i. \quad (4.13)$$

To find the location of the maximum, we again differentiate and set the derivative to 0 to obtain:

$$\frac{\partial \ell}{\partial \lambda} = \frac{n}{\lambda} - \sum_{i=1}^n x_i \quad (4.14)$$

and solving this equation in λ gives the mle for the exponential model:

$$\hat{\lambda} = \frac{n}{\sum_{i=1}^n x_i}. \quad (4.15)$$

4.3.3 Theoretical properties for the mle

The mle has the advantage of having numerous theoretical properties for any reasonable model. In particular, it has two notable properties known as asymptotic *consistency* and asymptotic normality. We say that these properties are *asymptotic* as they only hold exactly as the sample size goes to infinity, but we will usually consider that they hold approximately for a finite sample.

Consistency denotes the fact that although the mle is not unbiased, its bias vanishes as the sample size goes to infinity. The mle is always right when we collect an infinite

amount of data. This property can be seen as an analogue of the law of large number for the sample mean.

Normality denotes the fact that the mle follows an approximately normal distribution when the sample is large, with a standard error that can be computed. The exact expression of this standard error is beyond the course, but most software tools we make use of are able to report it.

4.4 Method of moments

The method of moments is an alternative strategy to obtain estimators for a general estimation problem. Although it is often simpler than the mle, it is somewhat inferior and its theoretical properties can be difficult to analyze. It is also difficult to apply in complex models where the parameter is no longer a number but a more general mathematical object.

The idea of the method of moment is to match the population and sample moments, and solve the set of equations to obtain an estimator. First, let us define what a *moment* is.

Suppose we have a sample X_1, \dots, X_n . We define the first population moment to be the population mean $\mu_1 = \mathbb{E} X$, and similarly, the first sample moment is the sample mean $M_1 = \frac{1}{n} \sum_{i=1}^n X_i$.

Now, for $j \geq 2$, we define the j^{th} centered population moment to be:

$$\mu_j = \mathbb{E}(X - \mathbb{E} X)^j, \quad (4.16)$$

and the j^{th} centered sample moment in a analogous fashion as:

$$M_j = \frac{1}{n} \sum_{i=1}^n (X_i - \bar{X})^j, \quad (4.17)$$

where \bar{X} is the sample mean. The second centered moment is usually called the variance, the third one the *skewness*, and the fourth one the *kurtosis*.

4.4.1 Example: exponential distribution

Let us illustrate the method of moments by a simple example, the exponential example we also considered for the mle. Suppose that X_1, \dots, X_n follow an exponential distribution with parameter λ . The population mean is then given by $\mu_1 = \lambda^{-1}$. The sample mean is simply given by $M_1 = n^{-1} \sum_{i=1}^n X_i$.

To obtain an estimator of λ according to the method of moments, we simply equate $\mu_1 = M_1$, and solve for the parameter λ . We thus obtain:

$$\frac{1}{\lambda} = \frac{1}{n} \sum_{i=1}^n X_i, \quad (4.18)$$

which we may solve to obtain the method of moments estimator:

$$\hat{\lambda} = \frac{n}{\sum_{i=1}^n X_i}. \quad (4.19)$$

4.4.2 Example: gamma distribution

Let us consider a slightly more complex example, the gamma distribution, with two parameters. Suppose that X_1, \dots, X_n follow a gamma distribution with shape k and scale θ . The p.d.f. is given by:

$$f_X(x) = \frac{1}{\Gamma(k)\theta^k} x^{k-1} e^{-x/\theta}, \quad (4.20)$$

and the mean and variance of the gamma distribution are given by

$$\mathbb{E} X = k\theta \text{ and } \sigma^2 = k\theta^2. \quad (4.21)$$

Note that applying the mle methodology manually here would be quite difficult, as the p.d.f. is complex, although it is still quite easy with the help of a computer. However, we may apply the method of moments to obtain the system of equations:

$$\begin{cases} M_1 = k\theta, \\ M_2 = k\theta^2, \end{cases} \quad (4.22)$$

which we may solve in the parameters to obtain two estimators $\hat{\theta} = M_2/M_1$ and $\hat{k} = M_1^2/M_2$. In this case, we needed to match two moments as we had two parameters. In general, we will have to match as many moments as there are parameters.

4.5 Uncertainty in estimation

In addition to providing a “guess” for the parameter, a statistician is also interested in quantifying the certainty or uncertainty of the given guess. Indeed, although observing 5 heads in 10 tosses, and 500 heads in 1000 tosses gives the same estimate $\hat{p} = 0.5$, having a larger sample size indicates that the second estimate is more confident.

As we have seen before, one possibility to do so could be to report the standard error of the estimate, which gives some indication as to the variability of the estimate. However, in some cases the standard error may fail to capture the whole picture, and we will prefer to report a range of values, called a confidence interval.

4.5.1 Confidence intervals

The idea of a *confidence interval* is to produce an estimate as a range of possible values instead of a single value. We would like this range of values to capture the likely possibilities of the parameter. As this range is produced from the data, the range itself is a random quantity, and thus can be analyzed probabilistically.

Let us define formally a $(1 - \alpha)$ confidence interval (e.g. for a 95% confidence interval, $\alpha = 0.05$). We may write a confidence interval as $[a(X), b(X)]$, where a is the lower bound, and b the upper bound, and the dependence on the data X has been made explicit. Then, we say that $[a(X), b(X)]$ is a $(1 - \alpha)$ confidence interval for θ if:

$$P_\theta(a(X) \leq \theta \leq b(X)) = 1 - \alpha, \quad (4.23)$$

where the probability is taken supposing that θ is the true value. We may interpret this as saying that if we repeat the experiment many times, and compute a 95% confidence interval in the same way every time, this interval will cover the true value of the parameter 95% of the experiments.

4.5.2 Confidence interval for a normal observation

We will now compute a confidence interval for a single normal observation. This is one of the most used forms of confidence intervals, as we have seen that mle are asymptotically normal, and hence this method may be used to produce confidence intervals for the mle.

Suppose that we observe a single observation $\hat{\theta} \sim \mathcal{N}(\theta, \sigma_\theta^2)$, where σ_θ is the standard error of our estimate (which is supposed known). We claim that the interval:

$$\left[\hat{\theta} - \sigma_\theta z_{1-\alpha/2}, \hat{\theta} + \sigma_\theta z_{1-\alpha/2} \right] \quad (4.24)$$

is a $(1 - \alpha)$ confidence interval, where we have defined $z_{1-\alpha/2}$ to verify, where Z is a standard normal random variable:

$$P(Z \leq z_{1-\alpha/2}) = 1 - \alpha/2. \quad (4.25)$$

Example 29 Suppose that after running an experiment, we obtain an estimate of $\hat{\theta} = 4.5$ with a standard error of 1.2. Let us compute a 95% confidence interval for θ . Using software, we may compute that $z_{1-\alpha/2} = 1.96$, and hence the interval is given by:

$$[4.5 - 1.96 \times 1.2, 4.5 + 1.96 \times 1.2] = [2.15, 6.85]. \quad (4.26)$$

4.5.3 Bootstrapping confidence intervals

The normal assumption is not always appropriate, and we will sometimes encounter cases in which we may not wish to use a normal approximation, or are not able to compute the standard error of the estimator. In these cases, the *bootstrap* is a general strategy that allows us to estimate a confidence interval from the data.

In general, we may compute a confidence interval by noting that if we knew a and b such that:

$$P(a \leq \hat{\theta} - \theta \leq b) = 1 - \alpha. \quad (4.27)$$

Indeed, this would directly imply (by algebraic manipulation) that:

$$P(\hat{\theta} - b \leq \theta \leq \hat{\theta} + a) = 1 - \alpha, \quad (4.28)$$

that is, the interval $[\hat{\theta} - b, \hat{\theta} + a]$ is a $(1 - \alpha)$ confidence interval for θ .

Unfortunately, we do not have access to the distribution of $\hat{\theta} - \theta$, as the distribution of the estimator may be complex. Instead, we propose to estimate this distribution. We estimate the distribution by artificially creating new datasets by resampling from our existing data.

We create a new dataset of the same size as our existing dataset by picking observations from our dataset at random (we may pick the same observation several times). We then

4 Estimation

compute the value of our estimator on this resampled dataset, and call its value θ^{boot} . By simulating this artificial dataset many times, we may estimate the distribution of $\theta^{\text{boot}} - \hat{\theta}$.

Now, the idea is that the distribution of $\theta^{\text{boot}} - \hat{\theta}$ is close to that of $\hat{\theta} - \theta$, and so instead of using a and b from the original distribution, we may have a and b be sample quantiles from the bootstrap distribution.

5 Regression

In the previous section, we have seen how to estimate parameters in the case that each observation was generated independently from the same parameter. However, we will commonly be interested in the case where we wish this parameter to depend on other covariates. For example, consider a drug for a cancer treatment. For a given person, we may wish to model the outcome of the drug as a Bernoulli random variable with probability of success p . Additionally, we may believe that that this probability of success p may depend on some other characteristics of the person under treatment, for example their age or the status of the cancer at the start of the treatment.

The idea of regression is to allow our parameter of interest to vary as a function of covariates, and estimate the relationship between the parameter of interest from the data. Mathematically, we may write $p = f(x)$, where x represents the covariate or covariates of interest. The goal of regression is then to estimate f , and understand how that may inform us about the relation between the covariates and the parameter.

5.1 Linear regression

Although it is possible to estimate the function f in full generality (usually called *non-parametric* regression, this requires substantial amounts of data to obtain a good estimate, and can be difficult to interpret. Instead, the most common model is that f is linear:

$$f(x) = \alpha + \beta x. \quad (5.1)$$

This allows us to reduce the problem of estimating a whole function to estimating some numbers: the intercept α , and the coefficient β . These parameters may now be estimated by the maximum likelihood strategy.

In general, we may be interested in several covariates, in which case we extend the function f as a linear function of several variables:

$$f(x_1, x_2) = \alpha + \beta_1 x_1 + \beta_2 x_2. \quad (5.2)$$

We may also use a maximum likelihood strategy to estimate any number of parameters, although it becomes progressively more analytically difficult.

5.1.1 Ordinary least squares

The most common case, and what is usually referred to as *ordinary least squares* or simply as linear regression, considers a normal observation model, where each observation

5 Regression

is normal with some given mean μ and variance σ^2 . We then suppose that the mean μ may depend on covariates in a linear fashion:

$$\mu = \alpha + \beta_1 x_1 + \beta_2 x_2 + \cdots \quad (5.3)$$

Additionally, we suppose that the variance σ^2 is constant.

Let us compute the estimate $\hat{\alpha}$, $\hat{\beta}$ in the case that there is only one covariate X . Suppose that we observe independent observations (y_i, x_i) where y_i is the outcome and x_i the covariate. We may write the joint likelihood as:

$$L(\alpha, \beta) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left\{-\frac{1}{2\sigma^2}(y_i - (\alpha + \beta x_i))^2\right\}. \quad (5.4)$$

The joint log-likelihood is then given by the log of the quantity:

$$\ell(\alpha, \beta) = -\frac{n}{2} \log(\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (y_i - (\alpha + \beta x_i))^2. \quad (5.5)$$

The first term is constant in the parameters, hence we may ignore it for the purpose of maximizing the likelihood in terms of the parameters. We thus wish to compute:

$$\arg \max_{\alpha, \beta} -\frac{1}{2\sigma^2} \sum_{i=1}^n (y_i - \alpha - \beta x_i)^2. \quad (5.6)$$

This is equivalent to minimizing the opposite:

$$\arg \min_{\alpha, \beta} \frac{1}{2} \sum_{i=1}^n (y_i - \alpha - \beta x_i)^2. \quad (5.7)$$

We may now compute the optimal value by differentiating and setting the derivatives to zero. Let us first consider the intercept α , we may compute the derivative:

$$\frac{\partial}{\partial \alpha} \frac{1}{2} \sum_{i=1}^n (y_i - \alpha - \beta x_i)^2 = -\sum_{i=1}^n (y_i - \alpha - \beta x_i) \quad (5.8)$$

Hence we have that the estimator $\hat{\alpha}$ is given by:

$$n\hat{\alpha} = \sum_{i=1}^n y_i - \left(\sum_{i=1}^n x_i\right)\hat{\beta}, \quad (5.9)$$

which is to say:

$$\hat{\alpha} = \bar{y} - \bar{x}\hat{\beta}. \quad (5.10)$$

5.1.2 Residuals

The ordinary least squares regression has a particularly geometrical interpretation (that other generalized linear models do not share). Indeed, we may think of the equation $y = \alpha + \beta x$ as the equation of a line, which we may interpret as the “line of best fit”. We may then be interested in the distance of the observed data points from the line.

The difference between the fitted value (the line) and the data point is residual, and is written:

$$r_i = y_i - \alpha - \beta x_i. \quad (5.11)$$

The size of the residuals can provide some attempt at measuring the goodness of fit of the line to the data, usually by the R^2 . Indeed, define the total sum of squares as $SS_{\text{tot}} = \sum_{i=1}^n (y_i - \bar{y})^2$. Also define the residual sum of squares as $SS_{\text{res}} = \sum_{i=1}^n r_i^2$. The coefficient of determination R^2 is then defined as:

$$1 - \frac{SS_{\text{res}}}{SS_{\text{tot}}}. \quad (5.12)$$

R^2 is a number between 0 and 1, representing the proportion of the variance explained by the regression. The R^2 is a simple summary of how well the model explains the data.

5.1.3 Generalized linear models and link functions

Instead of modelling a normal response, we may sometimes be interested in modelling some other response. Say we are interested in modelling the number of customers in a store, we may then believe that a Poisson model is most appropriate, and attempt to model the rate parameter λ . On the other hand, if we are modelling whether a treatment is successful or not, a Bernoulli model may be most appropriate, and attempt to model the success probability parameter p .

As before, we may write the parameter as a function of the covariates: e.g. the probability of success p of a treatment may depend on the patient’s age x , and we may write $p = f(x)$. As before, we may tempted to assume f to be a linear function and write $p = \alpha + \beta x$. However, we note that in this case, it may be possible to obtain $p < 0$ or $p > 1$! In order to alleviate this problem, we will consider a *link* function g , and model $g(p)$. To obtain the specific value for p , we will thus have:

$$g(p) = \alpha + \beta x, \quad (5.13)$$

which is equivalent to saying:

$$p = g^{-1}(\alpha + \beta x). \quad (5.14)$$

A common link for the Bernoulli case, called logistic regression, is the logistic function:

$$g(p) = \text{logit } p = \log \left(\frac{p}{1-p} \right). \quad (5.15)$$

It takes a value in $(0, 1)$ to $(-\infty, +\infty)$ and is a strictly increasing function.

5 Regression

A common link for the Poisson case is the logarithmic link:

$$g(\lambda) = \log \lambda. \quad (5.16)$$

As previously, the parameters may be estimated by maximum likelihood. However, there is no closed form analytical solution in most cases. Additionally, there is no natural geometric interpretation and notions such as residuals are more delicate.

5.2 Non-linear regression

In some cases, one may believe that the behaviour of the parameter is complex as a function of the covariates. In such cases, it may be more appropriate to consider a more general form of f , including methods that can learn the form of f from the data.

We will cover a popular method that attempts to make a reasonable trade-off between the complexity and the flexibility of the model: the generalized additive models (or *gam*). Mathematically, the *gam* models the parameter of interest θ with a link g that is known, and functions f_1, \dots, f_p that it estimates, where:

$$g(\theta) = \alpha + f_1(x_1) + f_2(x_2) + \dots + f_p(x_p). \quad (5.17)$$

As this method requires us to estimate entire functions f , it is somewhat more delicate to describe, although it is again inspired by the idea of maximum likelihood.

6 Hypothesis testing

A common question that statisticians face is that of deciding whether two quantities or groups are related. For example, one may wish to decide whether a drug was effective, that is, whether the proportion of people cured in the treatment group is higher than that in the control group. A researcher may also be interested in whether higher speeds lead to more fatal car crashes.

However, deciding these from a sample may sometimes be difficult given the randomness and uncertainty. Hypothesis testing studies the way in which we may quantify this uncertainty and produce decisions with theoretical guarantees.

6.1 Vocabulary of testing and first example

Suppose that we wish to test whether a drug for lowering blood pressure is effective. We thus observe the blood pressure of patients in a control group and in a treatment group. To set up the problem formally, we will consider a *null hypothesis*, which represents the “default” hypothesis, that we wish to reject. In this case, the null hypothesis would be that the average blood pressure is the same in each group. We will also need to consider an *alternative hypothesis*, which represents the hypothesis that we would like to validate. In this case, the alternative hypothesis would be that the average blood pressure is lower in the treatment group than in the control group.

Mathematically, we may write the blood pressure X_i as a random normal variable with mean μ_1 and variance σ^2 if it is from the control group, and with mean μ_2 and variance σ^2 if it is from the treatment group. We will thus write the two hypotheses:

$$H_0 : \mu_1 = \mu_2 \text{ v.s. } H_1 : \mu_1 > \mu_2. \quad (6.1)$$

In order to produce principled conclusions, we will wish to minimize our chance of error, although we will not be able to eliminate the possibility of error entirely. Indeed, the problems we consider are naturally random, and we can never logically eliminate the possibility that we are extremely unlucky, although we can limit the frequency of such outcomes by definition.

Mathematically, we will wish to control the rate of *type I* error, or false positive rate. We say we commit a false positive when the null hypothesis was true, but we instead declared it to be false. In statistical testing, the rate of type I error is called the *size* of the test, and is one of its most important quantity. As most of the applications of statistics occur in fields that tend to be conservative – for example scientists are much more worried about false discoveries than inconclusive experiments – guarding against the possibility of those false discoveries is important.

6 Hypothesis testing

The size of the test thus controls the frequency of these false discoveries. For example, if we consider a test with a size of 5%, we would expect about 1 in 20 discoveries to be a false discovery. If on the other hand we consider a test with a size of 1%, then we would expect about 1 in 100 discoveries to be a false discovery. When we discuss statistical tests, we will usually consider a family of tests, parametrised by the size or rate of false discovery.

On the other hand, it is also important to be able to detect differences when these exist. The ability of our test to do so is called the *power* of the test. More precisely, we say that we have a *type II* error, or false negative, when we fail to detect the alternative even though it is true. The complement of the rate of false negatives, i.e. the frequency at which our test is able to tell that the alternative is correct is called the power of the test.

The principle of statistical testing is to consider how unlikely it is that our data comes from the null model. For example, if the null model as above is that the two groups have the same average, we might consider, supposing that is true, what is the likelihood of obtaining a large difference in the sample means. Similarly, if we considered a null hypothesis that a coin is fair, but observed 8 heads out of 10, we might consider the probability to observe 8 or more heads in 10 tosses for a fair coin.

This probability is called the *p-value*, and characterises the result of our test. The ability to compute a p-value automatically allows us to produce a test of any size α . Indeed, rejecting the null hypothesis when $p < \alpha$ gives us a test of size α . However, there is no guarantee that the test has good power: although the test may not give us too many false positives, it may also not give true positives when required.

6.2 Permutation tests

Let us consider the example of testing for the blood pressure difference between the treatment and the control group. Suppose that we have $n = 20$ patients in each group, and have observations x_1, \dots, x_{20} in the treatment group, and y_1, \dots, y_{20} in the control group. A natural way to estimate the difference would be to consider the average sample difference $\hat{d} = \bar{y} - \bar{x}$. However, this quantity may be large simply by chance – although it is unlikely to be too large if there was no true difference. Can we characterise how large is large enough to rule out the possibility that it is due to chance?

An interesting observation is that, if there were no actual difference between the treatment and the control group, we could swap the observation from patient 10 in the treatment group with the observation from patient 5 in the control group, and that would be a stastically identical sample. Hence one possibility is to generate many fake samples by considering such permutations, and computing a value of \hat{d} for each of those permutations. This allows to characterise the distribution of \hat{d} , and compare with the value we have obtained, to compute a p-value.

More concretely, suppose that we observe a difference d_0 . We wish to understand how likely it is that we observe this difference or greater by simple chance if there were no actual difference, i.e. the p-value $P_{H_0}(\hat{d} > d_0)$. We may attempt to approximate this

value by simulating numerous permutations:

$$P_{H_0}(\hat{d} > d_0) \approx \frac{\text{number of permutations where } \hat{d} > d_0}{\text{total number of permutations}} \quad (6.2)$$

Permutation tests can sometimes be computationally expensive, as they require extensive simulations. However, they have the advantage of making minimal assumptions on the distribution of the measurements, and only suppose that the null hypothesis expresses that the two groups are statistically identical.

6.3 Pivot statistics

Another alternative to obtaining p-values for a test is to derive the distribution of \hat{d} (or a similar quantity) theoretically, and compare the observed value to the quantiles of that distribution. In general, analysing the theoretical distribution of these objects may be difficult.

For the example of testing differences between the two groups in terms of blood pressure, we may consider the t-statistic given by

$$t = \frac{\bar{x} - \bar{y}}{s_p \sqrt{2/n}}, \quad (6.3)$$

where we have

$$s_p = \sqrt{\frac{s_x^2 + s_y^2}{2}}, \quad (6.4)$$

with s_x^2 and s_y^2 being estimators of the variance of x and y respectively. Then, this t-statistic follows a t-distribution on $2n - 2$ degrees of freedom, and we may compare the computed value in the sample to the quantiles of the t-distribution.

A number of common statistical tests are defined by considering a test or *pivot* statistic, and deriving the distribution of the statistic under the null hypothesis.

6.4 Testing contingency tables

Contingency tables arise frequently, and we will often want to test for independence among such tables. For example, suppose that we have a trial with a treatment and control group, and recording whether the patient was cured or not at the end of the trial. In this case, it is natural to wonder whether the trial was effective, that is, whether there was some association between the treatment and the outcome. We may formulate this as a test with the null hypothesis being that the treatment and outcome are independent, and the alternative being that they are not.

We may record the values of the outcome in a contingency table, and describe the probabilities that a single observation lands in one of the cells as in table 6.1a. However, if we assume independence, then we have instead that the probability factorizes, and we

6 Hypothesis testing

	Treatment	Control
Cured	p_{11}	p_{12}
Not Cured	p_{21}	p_{22}

(a) General contingency table probabilities

	Treatment	Control
Cured	$p_{1\cdot}p_{\cdot 1}$	$p_{1\cdot}p_{\cdot 2}$
Not Cured	$p_{2\cdot}p_{\cdot 1}$	$p_{2\cdot}p_{\cdot 2}$

(b) Contingency table probabilities assuming independence

obtain table 6.1b, where we have written $p_{1\cdot}$ to be the probability that an observation is in the first row, and $p_{\cdot 1}$ the probability that it is in the first column.

We may then estimate those probabilities by simply counting the number of observations in the current table. For example, we may write:

$$\hat{p}_{1\cdot} = \frac{\text{\#observations in first row}}{\text{\#observations total}}, \quad (6.5)$$

and similarly for the other quantities. We may thus compute an expected number of observations in each cell, given by

$$E_{ij} = np_{i\cdot}p_{\cdot j} \quad (6.6)$$

where n is the total number of observations in the table.

Let O_{ij} be the actual count in cell (i, j) , we define the χ^2 statistic, given by:

$$\chi^2 = \frac{(O_{ij} - E_{ij})^2}{E_{ij}}. \quad (6.7)$$

Under the null hypothesis of independence, the χ^2 statistic has an approximate χ^2 distribution with 1 degree of freedom for the case of a 2×2 contingency table, the quantiles of which we may use to perform tests. This is commonly called the χ^2 *test of independence*. Because of the approximation, this test is only good if the cell counts are of moderate size (usually taken to be more than 5 or 10).

Alternatively, we may also consider our permutation strategy. In the permutation case, we may exchange randomly the outcomes of the patients under treatment and under control. In this specific case, it is possible to analyse theoretically the result of this permutation test, called *Fisher's exact test*. As this test is exact, it does not rely on any cell count size, and is preferred if we have cells with very low counts (less than 5).

7 Linear models

Linear (and generalized linear) models are some of the most ubiquitous models in statistics. They combine flexibility, power and ease of use, which makes them adapted for a wide variety of cases. In this chapter, we will consider how to use linear models to model common problems in statistics.

7.1 Coefficients of a linear model

7.1.1 Continuous variables in linear models

In a linear model, we model the mean response as a linear combination of the predictors. If y is our response, and x_1, \dots, x_p the predictors of interest, then we have that:

$$\mathbb{E} y = \alpha + \beta_1 x_1 + \dots + \beta_p x_p. \quad (7.1)$$

Most often, we will mostly be interested in the values of β , which describe how the mean response changes as a function of the covariates. For example, the value of β_1 describes the change in the mean response for each unit increase in x_1 , provided all other predictors $x_i, i \neq 1$ are held fixed. In addition, the effect of x_1 does not change as a function of the other covariates: a unit increase in x_1 always leads to an increase of β_1 in the mean response, no matter what the values of x_2, \dots, x_p are.

A special case of particular interest to us is when β_i is exactly zero, which represents the case that the covariate x_i has no influence on the average outcome when all other covariates are held fixed.

The interpretation of the coefficient as the effect of the covariate controlling for all other covariates present in the model is particularly important. For example, suppose that we consider a response that is only linked to the height of a person. Then, we would hope that the coefficient for the regression on the weight is (close to) zero, and that would indeed be the case when we have the height as a covariate. However, if height is not included as a covariate, the regression will attempt to “explain” the variation by using the weight, as height and weight are linked.

7.1.2 Categorical variables in linear models

In addition to numerical random variables, we will often wish to model categorical random variables in regressions. However, as categorical variables are not numerical, we cannot use the same formula directly. Instead, categorical variables are modelled using *dummy coding* in linear regression.

7 Linear models

Suppose that we have a categorical variable x with three levels, named A , B and C . How would we represent a linear regression of y on x ? The standard convention is to consider the value of x and consider all possible levels:

$$y = \alpha + \begin{cases} \beta_A & \text{if } x = A, \\ \beta_B & \text{if } x = B, \\ \beta_C & \text{if } x = C. \end{cases} \quad (7.2)$$

This specific representation has too many variables, and is non-identifiable. Indeed, we could add some arbitrary quantity to α , and take it away from all the β to obtain the same result. We will thus assume by convention that $\beta_A = 0$. Hence the usual representation for a single variable is given by:

$$y = \alpha + \begin{cases} 0 & \text{if } x = A, \\ \beta_B & \text{if } x = B, \\ \beta_C & \text{if } x = C. \end{cases} \quad (7.3)$$

We may generalize this to a regression on several variables, some of which may be categorical. For example, suppose that x_1 is categorical, and x_2 and x_3 are numerical. Then the linear regression is given by:

$$y = \alpha + \beta_2 x_2 + \beta_3 x_3 + \begin{cases} 0 & \text{if } x = A, \\ \beta_B & \text{if } x = B, \\ \beta_C & \text{if } x = C. \end{cases} \quad (7.4)$$

We may interpret the coefficient of the category as the difference in mean response from that category to the first category, which may be considered a reference category.

7.1.3 Generalized linear models

In generalized linear models, it is no longer the mean response that is modelled as a linear combination of the covariates, but rather some function of it (usually called the link function). This not only changes the value of the mean response for a given set of covariates, but also how the covariates affect the value of the mean response.

More concretely, let us write a generalized linear model with link g :

$$g(\mathbb{E} y) = \alpha + \beta_1 x_1 + \cdots + \beta_p x_p. \quad (7.5)$$

We may then interpret β_1 as the increase in $g(\mathbb{E} y)$ for each unit increase in x_1 , all other covariates being fixed. However, if we wish to understand the effect of x_1 on the mean response $\mathbb{E} y$, we need to invert the function g .

Poisson regression Let us consider the case of a Poisson regression. In the Poisson regression, the log link $g = \log$ is commonly used, and we denote the mean response λ . The Poisson regression then corresponds to the equation:

$$\log \lambda = \alpha + \beta_1 x_1 + \cdots + \beta_p x_p. \quad (7.6)$$

We may exponentiate both sides to obtain:

$$\lambda = e^\alpha \cdot e^{\beta_1 x_1} \cdots e^{\beta_p x_p}. \quad (7.7)$$

Hence, if we have a unit increase in x_1 , the new value of λ' is given by

$$\lambda' = e^\alpha \cdot e^{\beta_1(x_1+1)} \cdots e^{\beta_p x_p} = \lambda e^{\beta_1}. \quad (7.8)$$

Hence we may say that for each unit increase in x_1 , we see a multiplicative increase of e^{β_1} in λ .

Logistic regression In the case of logistic regression, we have a similar property. The link function is now given by the logit function, defined as $g(p) = \log(p/(1-p))$. Let $\pi = p/(1-p)$ be the *odds ratio*, we may then write the link function as $g(p) = \log \pi$.

Hence the logistic regression may be defined as:

$$\log \pi = \alpha + \beta_1 x_1 + \cdots + \beta_p x_p. \quad (7.9)$$

Similarly to the Poisson regression, we may interpret the coefficients in this case as characterising the multiplicative change in the odds-ratio π for each unit increase in the covariate. Indeed, by exponentiating both sides of eq. (7.9), we obtain the expression for the odds-ratio:

$$\pi = e^\alpha \cdot e^{\beta_1 x_1} \cdots e^{\beta_p x_p}. \quad (7.10)$$

Hence we may interpret e^{β_1} as the multiplicative change in the odds-ratio for every unit increase in x_1 .

7.2 Significance in a linear model

7.2.1 Testing significance of a single coefficient

A common and important question in statistics is understanding whether a given covariate affects the outcome, once other covariates have been controlled for. For example, one may wonder if smoking affects the probability having lung cancer, controlling for socio-economic status. Or we may be interested in understanding whether gender affects salary, controlling for education level, experience etc.

Linear modelling provides a simple yet powerful framework to model and answer these questions. We can control for a variable by including it in the regression, and the question of whether the covariate of interest has an impact on the outcome can be simply stated as whether the coefficient of that variable is zero.

We may thus be interested in the following hypothesis test of $H_0 : \beta_1 = 0$ v.s. the alternative $H_1 : \beta_1 \neq 0$. Most statistical software automatically performs this test, and computes the p-value. In addition, R provides a visual representation of the p-value to indicate which variables are significant.

7.2.2 Confidence interval for a single coefficient

Closely linked to the notion of significance of a coefficient, we may also compute the standard error of a single coefficient. This allows us to compute confidence intervals for the coefficient of interest. In the case of the linear model with normal errors, the estimates of the coefficients follow an exact normal distribution, and in other cases they follow an approximate normal distribution.

We can thus follow the procedure in section 4.5.2 to produce a confidence interval for a single coefficient from the estimate and the standard error. Indeed, let $\hat{\beta}_1$ be our estimate for β_1 , and σ the standard error of the estimate, a $1 - \alpha$ confidence interval for β_1 is then given by:

$$[\hat{\beta} - \sigma z_{1-\alpha/2}, \hat{\beta} + \sigma z_{1-\alpha/2}], \quad (7.11)$$

where $z_{1-\alpha/2}$ denotes the quantiles of a standard normal, defined by $P(Z \leq z_{1-\alpha/2}) = 1 - \alpha/2$ for Z a standard normal variable.

7.3 Model selection

Model selection denotes the problem of selecting the “best” linear model among a set of possible candidate models. Indeed, suppose that we have a number of predictors x_1, \dots, x_p , and wish to select the best and discard those that have little predictivity or influence on the result.

7.3.1 Information criterion

A common strategy to select such a best model is to consider some measure of how good a model is, and then pick the model (among those under consideration) that maximizes this measure. For example, suppose that we wish to consider two possible models, $y = \alpha + \beta_1 x_1$ and $y = \alpha + \beta_1 x_1 + \beta_2 x_2$, can we decide (by only looking at the data) which of the two is the better one?

A possibility would be to consider the coefficient of determination, or R^2 . However, although the R^2 is a measure of how well the model approximates the data at hand, it has the undesirable property of being always increasing. Indeed, as the model gets larger and has more variables, it is always possible to reduce the on our given dataset. However, larger models also have larger variance, and may not always correspond to the better solution.

Instead, we wish to define a mixed criterion, that not only considers how well the model fits the data, but also how complex the model is. In the context of linear models, there are three commonly used criteria: the *AIC*, or Akaike’s information criterion, the *BIC*, or Bayesian information criterion, and the adjusted R^2 .

Let n be the number of observations, k be the number of variables in the model, and RSS be the residual sum of squares. Then the above criteria are defined as follows:

$$\begin{aligned} \text{AIC} &= 2k - 2 \log \text{RSS}, \\ \text{BIC} &= k \log n - 2 \log \text{RSS}, \\ \text{adj} - R^2 &= 1 - \frac{\text{RSS}/(n - k)}{\text{TSS}/(n - 1)}. \end{aligned}$$

For the AIC and BIC, lower is considered better, whereas for the adjusted R^2 , higher is considered better.

These criteria measure the tradeoff between model complexity and improved fit of the model. The AIC and BIC have some theoretical underpinning, whereas the adjusted R^2 is defined to mimic the R^2 , but unlike the latter, does not necessarily increase as model complexity increases.

In the context of estimation, the BIC is one of the most common criterion used to select the model, as it tends to select models that have fewer variables, which are easier to interpret. On the other hand, the AIC tends to select models with more variables, which may have better predictivity, especially when the true model is not among the model that are being selected.

7.3.2 Which models to select from?

A parallel problem to that of determining how to select the best model, we also need to produce a list of models to choose from. In some cases, there may be candidate models that we wish to compare manually. However, in most cases, we will only have access to a list of variables that may be of interest. In this case, a common strategy is to consider *best subsets* selection.

In the case of best subsets selection, we algorithmically consider all models that are possible from the given variables of interest. For example, if we had only two covariates x_1 and x_2 , we could form the following four different models:

$$\begin{aligned} \mathbb{E} y &= \alpha, \\ \mathbb{E} y &= \alpha + \beta_1 x_1, \\ \mathbb{E} y &= \alpha + \beta_2 x_2, \\ \mathbb{E} y &= \alpha + \beta_1 x_1 + \beta_2 x_2. \end{aligned}$$

We may then use any of the criterion from the previous section to choose the model that balances the best between corresponding to the observed data and being parsimonious.

7.4 Penalized linear models

The standard OLS estimator is unbiased for normal errors. However, we have seen that this is not always a desirable property, and we can in fact often do better by introducing some amount of bias in order to reduce the bias. This idea is manifested through the

idea of penalization in the context of regression, where we artificially move the estimates of the coefficients closer towards zero. When done carefully, this can make our estimates more precise by having a more favorable bias-variance tradeoff.

7.4.1 Ridge regression

In order to penalize the parameters, we may modify the problem slightly by penalizing the objective function that defines the estimates. Recall that for OLS (with a single variable), the estimate $\hat{\alpha}^{OLS}$, $\hat{\beta}^{OLS}$ minimize the objective:

$$\text{RSS} = \sum_{i=1}^n (y_i - \alpha - \beta x_i)^2 \quad (7.12)$$

The *ridge regression* estimates, on the other hand, minimize the penalized objective

$$\sum_{i=1}^n (y_i - \alpha - \beta x_i)^2 + \lambda \beta^2, \quad (7.13)$$

where $\lambda > 0$ is some tuning parameter that controls the bias-variance tradeoff.

Note that for $\lambda = 0$, we recover the simple ordinary least squares estimator which is unbiased but has potentially high variance, whereas as $\lambda \rightarrow \infty$, we set $\hat{\beta} = 0$, which has zero variance but potentially high bias.

The choice of λ is a delicate and important part of obtaining good estimates using ridge regression. However, there exists a completely automated procedure, called *cross-validation*, that can automatically select the value of λ from the data.

In general, ridge regression improves the accuracy of the estimates but remains mostly similar to OLS. It has the advantage of being able to compensate for collinear estimators.

7.4.2 Lasso regression

A much more recent technique, lasso regression combines the idea of penalization and that of model selection to create a technique that can automatically select a model from the data and estimate the coefficients at the same time.

Although following a similar principle as the ridge regression by penalizing the objective function, the lasso objective function has some interesting mathematical properties that allows it to function as model selection tool. The Lasso estimator minimizes the following function:

$$\sum_{i=1}^n (y_i - \alpha - \beta x_i)^2 + \lambda |\beta|, \quad (7.14)$$

where again $\lambda > 0$ is a tuning parameter. Due to the mathematical properties of the absolute value function $|\beta|$, this specific penalization will sometimes give the exact estimate $\hat{\beta} = 0$, which can be interpreted as not selecting the corresponding variable.

The Lasso is a particularly useful variable selection tool in cases where the number of predictors is greater than the number of observations. These problems often arise in modern datasets such as genetics or other signals. Indeed, OLS cannot fit a model

where the number of predictors is greater than the number of observations, so the usual techniques such as best subset selection cannot be applied. On the other hand, Lasso is a fast and efficient method that can be used even with millions of variables but only thousands of observations.

8 Bayesian statistics

Bayesian statistics is a somewhat different paradigm of statistics, in which one considers the parameters not as unknown fixed values, but instead as random values. Throughout these notes, we have considered a parameter of interest (e.g. the probability of a coin landing on head), as a fixed value (i.e. a number) that is unknown to us. On the other hand, the Bayesian paradigm models this parameter as a random number, with a distribution that models our belief about the parameter.

8.1 Priors

This initial distribution placed on the parameter of interest θ is called the *prior*, and is usually denoted by $\pi(\theta)$. This is often viewed as a distribution that captures our *a priori* or existing knowledge about the parameter. For example, we may attempt to encode the fact that a coin is likely to not be too biased, i.e. that $p \approx 0.5$ in the form of a prior.

The choice of a specific prior can sometimes be a contentious issue, although in practice most practitioners attempt to encode some notion of uncertainty by using a prior with large variance.

8.2 Likelihood and posterior

Given a prior $\pi(\theta)$, we would now like to understand how the observed data X should change our belief about the distribution of the parameter θ . By Bayes' formula, we may write the *posterior* distribution of θ given X as:

$$\pi(\theta | X) = \frac{f(x | \theta)\pi(\theta)}{f(x)} \quad (8.1)$$

where $f(x | \theta)$ denotes the distribution of the data for θ fixed, in other words the likelihood of the data.

We note that as the observation X is usually considered as fixed, we may write the posterior up to a constant of proportionality as:

$$\pi(\theta | X) \propto L(\theta | X)\pi(\theta), \quad (8.2)$$

where we have written the likelihood $L(\theta | X) = f(x | \theta)$. The constant of proportionality may be recovered by integrating θ over the entire range as a probability distribution must integrate to 1.

8.3 Bayesian estimators

The Bayesian paradigm allows us to obtain the posterior distribution of our parameter (at least in principle). However, we would usually wish to obtain a single number representing our estimate, instead of an entire distribution that may be difficult to understand. The most common Bayesian estimator is the posterior mean, that is, the mean of the parameter under the posterior distribution:

$$\hat{\theta}^B = \mathbb{E}_{\theta|X} \theta = \int_{-\infty}^{+\infty} \theta \pi(\theta | X) d\theta. \quad (8.3)$$

Bayesian estimators always have some bias. However, the prior also acts as a regularizing mechanism, which implies that Bayesian estimators often have lower variance than their frequentist counterparts. In many cases, this gives Bayesian estimators many desirable properties, and they tend to perform well when there is not a lot of data or the data is not very informative.

8.4 Examples

In most cases, characterising the posterior distribution of θ can be difficult, and we must resort to numerical methods such as MCMC (Markov Chain Monte Carlo). However, in some special cases, the posterior will be of the same family of the prior. In this case, we are often able to solve analytically for the posterior parameters, and we say the prior is *conjugate*.

8.4.1 Binomial model

Let us first consider a binomial model. Suppose that $X | \text{Binom}(n, p)$, with n a known number, and p following a $\text{Beta}(\alpha, \beta)$ distribution, for some $\alpha, \beta > 0$ known. The p.d.f. of p is given by:

$$\pi(p) = \frac{p^{\alpha-1}(1-p)^{\beta-1}}{B(\alpha, \beta)} \quad (8.4)$$

where $B(\alpha, \beta)$ is a normalizing constant.

Let us compute the posterior distribution $\pi(p | X)$ of p . We will work up to a constant (in p) of proportionality, as this will simplify our computation substantially. We have that the posterior is given by:

$$\begin{aligned} \pi(p | x) &\propto L(p | x) \pi(p) \\ &\propto \binom{n}{x} p^x (1-p)^{n-x} p^{\alpha-1} (1-p)^{\beta-1} \\ &\propto p^{x+\alpha-1} (1-p)^{n-x+\beta-1}, \end{aligned}$$

where we have used the fact that $\binom{n}{x}$ is constant in p . We may now recognise that $\pi(p | x)$ is again proportional to a Beta distribution with parameters $\alpha + x, \beta + n - x$, hence we

have that $p \mid x \sim \text{Beta}(\alpha + x, \beta + n - x)$. The mean of a $\text{Beta}(\alpha, \beta)$ distribution is given by $\alpha/(\alpha + \beta)$, hence the posterior mean of p is given by:

$$\hat{p}^B = \frac{\alpha + x}{\alpha + \beta + n}. \quad (8.5)$$

Note that this corresponds to the frequentist estimator x/n in the case that we have “seen” a total of $\alpha + x$ heads and $\beta + (n - x)$ tails. The parameters α and β are thus often interpreted as *prior counts*. On the other hand, the estimator \hat{p}^B has expectation $\mathbb{E} \hat{p}^B = (\alpha + np)/(\alpha + \beta + n)$, and so is not unbiased.

8.4.2 Normal model

Let us consider a second example of a normal model. Let $X \sim \mathcal{N}(\mu, 1)$, and let $\mu \sim \mathcal{N}(0, 1)$. Then, we have that

$$\pi(\mu) = \frac{1}{\sqrt{2\pi}} e^{-\frac{1}{2}\mu^2}. \quad (8.6)$$

By Bayes’ rule, we may compute the posterior distribution of μ up to a constant of proportionality with:

$$\begin{aligned} \pi(\mu \mid x) &\propto L(\mu \mid X) \pi(\mu) \\ &\propto \exp\left\{-\frac{1}{2}(x - \mu)^2\right\} \exp\left\{-\frac{1}{2}\mu^2\right\} \\ &\propto \exp\left\{-\frac{1}{2}(x^2 - 2\mu x + 2\mu^2)\right\} \\ &\propto \exp\left\{-\frac{1}{2}(2\mu^2 - 2\mu x)\right\} \\ &\propto \exp\left\{-\frac{1}{2}[2(\mu^2 - x/2)^2 - x^2/2]\right\} \\ &\propto \exp\left\{-(\mu^2 - x/2)^2\right\} \end{aligned}$$

Hence we see that the posterior density is proportional to a normal density with mean $X/2$ and variance $1/2$. Hence the posterior distribution of μ is given as $\mu \mid X \sim \mathcal{N}(X/2, 1/2)$, and we have that the posterior estimator is $\hat{\mu}^B = X/2$.

As we can see again, this is not an unbiased estimator, as we have that $\mathbb{E} \hat{\mu}^B = \mu/2$, but it introduces some penalization by moving the estimate towards zero.

9 Prediction

The problem of prediction has returned as one of the most important and interesting problems in statistics. With the advent of massive computational power and data collection capacities, modern prediction algorithms can accurately estimate complex models.

9.1 Defining prediction

In statistical prediction, we are not so much interested in understanding the true latent state of nature, as much as predicting observable outcomes for new individuals. For example, classical estimation may be interested in understanding how an individual's diet affects their chance of heart disease, whereas a prediction take on this problem would be to guess whether that person would have heart disease based on their diet.

Although the difference may seem subtle, it is important, as prediction is in general less sensitive to confounding: the *why* is not so important as the *what*. Consequently, we should apply predictive methods with care when we desire to understand the process behind how the predictions were achieved. In particular, they can be delicate to apply in scientific contexts, where we may be interested in understanding the “true reason” behind some phenomenon.

9.1.1 What is a good prediction

Just like in the context of estimation, we wish to create a quantity to measure the goodness of our predictions. We usually call this quantity the *loss*, which we may think of as a proxy to the amount of error that is committed.

As in the context of estimation, the most common loss in use is the *square loss*, which is defined as:

$$L(y, \hat{y}) = (y - \hat{y})^2. \quad (9.1)$$

However, note that unlike the square loss for estimation, this loss is interested in the error we commit in the prediction of the outcome. In particular, it is possible to estimate this quantity (by comparing the prediction to the true value), whereas the estimation error cannot be estimated directly (as indeed we do not know the true value of the parameter). The ability to estimate this error is central in most of the predictive techniques.

9.1.2 Overfitting and underfitting

As we saw in the previous chapters, the problem of estimation is characterised by a bias-variance tradeoff, where more complex models would be better as the amount of data increased, but could perform poorly when the amount of available data is limited.

Models in prediction face a similar tradeoff, usually called the overfitting problem. Indeed, in building a model to predict a dataset, the training data we use may have some idiosyncracies due its randomness. Although we wish to capture the potentially complex pattern in the available, we must take great care to not also capture coincidences that may appear in the data.

Overfitting often manifests itself by having a training error (that is, the average loss on the training set) being much lower than the test error (that is, the average loss on an independent test set).

9.2 Linear models for prediction

The ubiquitous linear models we have used throughout can indeed also be used to predict outcomes. Indeed, suppose that we have fit a linear model with estimated coefficients $\hat{\alpha}$ and $\hat{\beta}$. Then, we may predict the outcome by simply plugging the value of the coefficients into the defining equation of the model:

$$\hat{y} = \hat{\alpha} + \hat{\beta}x \quad (9.2)$$

Example 30 We have analysed a dataset of diamond prices containing the price for which the diamond was sold and some characteristics of the diamond.

We model the price y of the diamond as a linear function of the weight x of the diamond (in carats). The estimated coefficients are $\hat{\alpha} = -2256$ and $\hat{\beta} = 7756$. We deduce that the predicted price for a diamond weighing 2 carat is given by:

$$\hat{y} = -2256 + 7756 \times 2 = 13256. \quad (9.3)$$

It may in some cases be preferable to transform the outcome of interest. Suppose instead that we regressed $\log y$ as a linear function of the weight x of the diamond. We obtain the following estimates $\hat{\alpha} = 6.215$ and $\hat{\beta} = 1.970$. We deduce that the predicted price for a diamond weighing 2 carat is given by;

$$\log \hat{y} = 6.215 + 1.97 \times 2 = 10.155 \quad (9.4)$$

from which we may deduce that $\hat{y} = 25719$.

9.2.1 Confounding in prediction

Confounding can be a serious issue in estimation, as it may cause us to obtain incorrect estimates when we did not account for all the possible covariates. However, in the context of prediction, confounding is less important. Although confounded predictors cannot be

interpreted in a causal fashion, they may be useful nonetheless in predicting the value of the outcome.

For example, in a dataset of diamond prices, we regress the price of the diamond on its colour rating. However, we see that on average, the better the colour rating is, the cheaper the diamond becomes. We have seen that from a regression point of view, this is due to a confounding effect of the weight of the diamond.

Nevertheless, if we were only able to know the colour rating (and not the weight), it would be correct to predict better ratings as cheaper. Indeed, the average diamond that we see with a better colour rating will be smaller, and hence cheaper on average.

9.2.2 Uncertainty in prediction

Just as there is uncertainty in our estimation of our coefficients, there is natural uncertainty in the prediction we can give for a given observation. We can characterise this uncertainty by giving a *prediction interval*, which describes a range of values that will likely contain the true observed value.

In general, this interval reflects two types of uncertainty: one, due to the uncertainty in estimating the model – e.g. the uncertainty in the value of the coefficients in the case of a linear model. As we accululate more data, we would hope that this uncertainty diminishes.

However, the prediction interval must also reflect the uncertainty due to the inherent randomness of the model. For example, consider the problem of predicting the number of radioactive decay events within a given time frame. Even if we knew the true average rate of those events, the actual realised number has inherent random variation that cannot be perfectly predicted.

9.3 Classification

We say a prediction problem is a *classification* problem if the target to be predicted is some categorical variable: usually some success or failure variable, or some target group. In the context of classification, the usual loss is given by a 0 – 1 loss, which measures the average percentage of correct predictions.

9.3.1 Performance of classification

The 0 – 1 loss is a good measure of loss whenever the groups are reasonably balanced, that is, in the population, the proportion of success outcomes is similar to the proportion of failure outcomes. If this is not the case, then the average percentage of correct predictions can be a misleading measure.

For example, consider attempting to predict credit card defaults. The rate at which credit card default occurs is extremely low, let us say less than 5% of the time. Then if we always predict no default, our accuracy would be about 95%, as indeed, 95% of the time there is no default. On the other hand, we are clearly not capturing any information in this prediction, and so we would like a quantity that can measure this fact.

9 Prediction

We can consider in this case the entire contingency table, which may describe using the following four rates. First, the rate of true positives, which is the proportion of the time we predict a success given that the true result was indeed a success. The rate of true negatives, which is the proportion of the time we predict a failure given that the true result was a failure. These former two correspond to correct predictions.

On the other hand, we can also consider the rate of false positives, which is the proportion of time we predict a success when the true value is a failure, and the rate of false negatives, which is the proportion of time we predict a failure when the true value is a success.

By definition, we always have that the rate of true negatives and false positives sum to 1, and that the rate of true positives and false negatives sum to 1. It is thus customary to consider the following two measures instead.

The *precision* of a decision method is given by the number of true positives over the total number of positive claims, whereas the *recall* of a decision method is given by the number of true positives over the true number of positive claims. Both quantities are between 0 and 1, and higher values indicate better performance.

We can also combine these two values into the F_1 score, which is the harmonic mean of the precision and the recall, given by:

$$F_1 = \frac{2pr}{p+r} \quad (9.5)$$

where p is the precision and r is the recall.

For most methods, we can trade off between the precision and the recall. Indeed, by being more conservative in claiming a positive outcome, we will reduce our total number of positive claims, and thus usually improve our precision. On the other hand, by being less conservative, and claiming more positive outcomes, we will increase the number of true positives, which improves our recall. This tradeoff may be captured in a precision-recall curve.

9.3.2 Linear models for classification

Logistic regression provides a way to use linear regression as a classifier. Indeed, logistic regression regresses the log-odds ratio of success as a linear function of the predictors, as in the equation below:

$$\log \frac{p}{1-p} = \alpha + \beta_1 x_1 + \cdots + \beta_p x_p. \quad (9.6)$$

We may thus use this equation to predict the probability of success for a given observation, by writing:

$$\hat{p} = \frac{e^{\hat{\alpha} + \hat{\beta}_1 x_1 + \cdots + \hat{\beta}_p x_p}}{1 + e^{\hat{\alpha} + \hat{\beta}_1 x_1 + \cdots + \hat{\beta}_p x_p}}. \quad (9.7)$$

Now, we may choose a cutoff (say $p = 0.5$), above which we will classify the outcome as a success, and below which we classify the outcome as a failure.

Now, as we vary the cutoff value p , we may trade off between the precision and the recall. As we increase the value of the cutoff towards 1, we only classify as a success those observations we are very certain will be successes.

9.3.3 Linear discriminant analysis

Linear discriminant analysis is a linear method that is similar to logistic regression for classification, although it can easily handle more than two classes (success/failure). It is a method that is somewhat inspired by a Bayesian idea.

Suppose that we have observations (y_i, x_i) , where y_i is the class of the observation, and x_i is some covariate. Suppose that we have l classes, numbered from 1 to l . We assume that for a given class k , the distribution of a covariate is given by:

$$x \mid y = k \sim \mathcal{N}(\mu_k, \sigma^2). \quad (9.8)$$

We thus assume that for a given class, the covariates follow a normal distribution with a mean corresponding to that class.

We now wish to predict the class of an observation given the covariate. We can thus compute by Bayes rule:

$$P(y = k \mid x) \propto \frac{f_X(x \mid y = k)\pi(k)}{1}, \quad (9.9)$$

where $\pi(k)$ is a prior on the class. Plugging the normal density in, we have that:

$$\begin{aligned} \log P(y = k \mid x) &\propto \exp\left\{-\frac{1}{2\sigma^2}(x - \mu_k)^2\right\}\pi(k) \\ &\propto \exp\left\{-\frac{1}{2\sigma^2}(\mu_k^2 - 2x\mu_k)\right\}\pi(k) \end{aligned}$$

We may naturally predict the class as the class with the highest probability. We will thus compute the log-probability for each class, also called the discriminant:

$$\delta_k(x) = \frac{x\mu_k}{\sigma^2} - \frac{\mu_k^2}{2\sigma^2} + \log \pi_k \quad (9.10)$$

To be able to compute the above, we need to estimate the prior probabilities π_k , the group means μ_k , and the pooled variance σ^2 .

To do so, we simply plug-in the values we observe from the data. Let n_k be the number of observations in group k , and n be the total number of observations. In particular, we compute the prior class probabilities and the group means from the data:

$$\begin{aligned} \hat{\pi}_k &= \frac{n_k}{n}, \\ \hat{\mu}_k &= \frac{1}{n_k} \sum_{i:y_i=k} x_i, \\ \hat{\sigma}^2 &= \frac{1}{n} \sum_{k=1}^l \sum_{i:y_i=k} (x_i - \mu_k)^2. \end{aligned}$$

We may extend the linear discriminant analysis to several covariates x_1, \dots, x_p by positing a multivariate normal distribution. Linear discriminant analysis can be a useful method to obtain a simple classification rule for continuous data.

9.3.4 Naive Bayes

Naive Bayes is another empirical Bayesian paradigm for classification, that is better adapted for large numbers of binary predictors. Suppose that we have observations $(y^{(i)}, x_1^{(i)}, \dots, x_p^{(i)})$, where x_1, \dots, x_p are binary predictors, and y_i is a 0 – 1 indicator of the class.

By Bayes' rule, we may obtain the posterior class probability by:

$$P(Y = 1 \mid X_1 = x_1, \dots, X_p = x_p) \propto P(X_1 = x_1, \dots, X_p = x_p \mid Y = 1) P(Y = 1). \quad (9.11)$$

However, estimating all the joint conditional probabilities for $P(X_1 = x_1, \dots, X_p = x_p \mid Y = 1)$ requires estimating $2^p - 1$ values, which would be difficult. We thus make the “naive” assumption, which is that the predictors are independent conditional on the class. In particular, this gives that:

$$P(X_1 = x_1, \dots, X_p = x_p \mid Y = 1) = P(X_1 = x_1 \mid Y = 1) \cdots P(X_p = x_p \mid Y = 1). \quad (9.12)$$

We may now estimate each probability from the data.

$$P(X_1 = x_1 \mid Y = 1) = \frac{\text{number of observations } X_1 = x_1 \text{ and } Y = 1}{\text{total number of observations with } Y = 1}. \quad (9.13)$$

By combining all such estimates, we may obtain the estimated probability by eq. (9.12).

The Naive Bayes estimator is particularly adapted to large amounts of categorical data (in particular binary data), and is historically a good model for text classification such as spam classification. In this context, each variable X_i is an indicator of whether a given word is in the document. For example, X_1 might be 0 if the email does not contain the word “business”, and 1 otherwise. X_2 might be the same for the word “meeting”.

Laplace smoothing A particular problem can arise when one of the estimated probabilities is exactly 0. Indeed, as we are combining the probabilities by multiplying them, if one of the exact probabilities is exactly 0, then the total estimated probability is 0. However, this can be undesirable, as it ignores the other evidence we have.

To solve this problem, we may add a small number to both the numerator and the denominator of the estimate, called the Laplace smoothing, for example:

$$P(X_1 = x_1 \mid Y = 1) = \frac{\text{number of observations } X_1 = x_1 \text{ and } Y = 1 + 0.5}{\text{total number of observations with } Y = 1 + 1}. \quad (9.14)$$

This can also be seen from a Bayesian perspective, by assuming that we have placed a prior on each probability (for example, see section 8.4.1).

10 Machine learning

With the increase in computational power and data collected in the world today, predictive methods that are able to make use of large amounts of data in an effective fashion have seen an increase in popularity.

In this section, we discuss some models that can be very flexible, and are able to generate impressive predictions when given large amounts of data. On the other hand, these models are often difficult to understand and interpret, and require a reasonable amount of data to perform.

10.1 Tree-based methods

Trees are a natural way to encode potentially complex decisions. They encode decision rules as a sequence of dichotomies on a given variable at a time.

10.1.1 Decision trees

We first consider *decision trees*, which encode a categorical outcome as a sequence of dichotomies. For example, consider fig. 10.1, a decision tree that encodes the species of an iris tree according to some covariates.

Each node of the tree represents a decision, which is a dichotomy on the value of a given variable. The values that are less than the given threshold are classified in one branch, and the values that are greater than are classified in another branch.

Each leaf represents a distribution of the categories in that node. That is, it considers all observations that verify the dichotomies corresponding to the path from the root to that leaf, and plots their distribution.

To use such a tree to classify a given observation, we may simply follow the observation down the tree according to the nodes, and pick the most likely class in that subset.

10.1.2 Regression trees

Trees may be used in a similar fashion to predict continuous outcomes, and are usually called *regression trees* in this context. Similarly to a decision tree, regression trees correspond to a sequence of dichotomies. However, instead of predicting the most likely class at each leaf node, the regression tree instead predicts the average value of the outcome at that node. For example, we may consider a regression tree to predict the log psa in prostate data, as in fig. 10.2.

Regression trees produce a piecewise constant regression, as within each leaf node the predicted value is constant. However, they can model fairly complex interactions, and

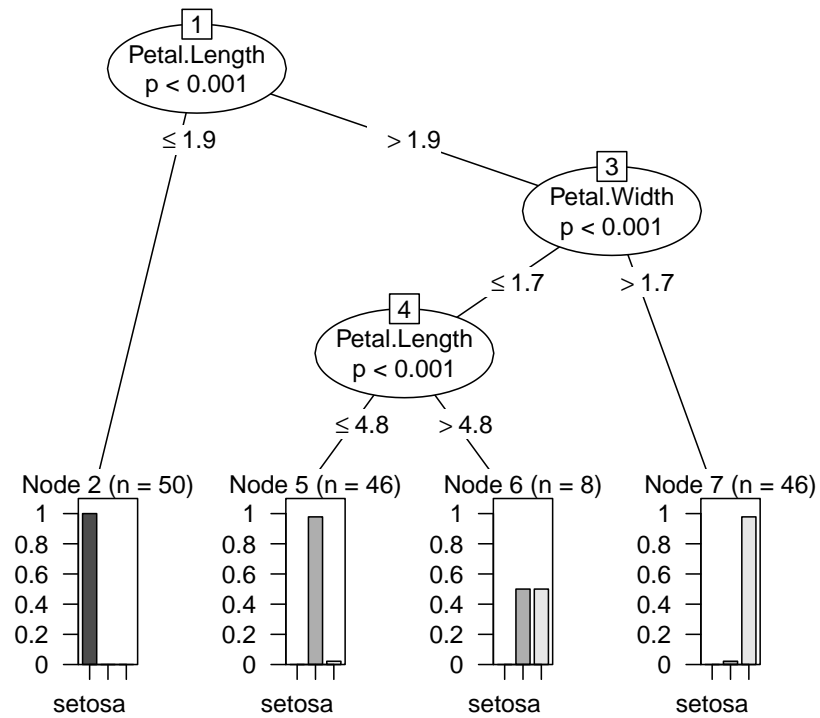


Figure 10.1: Decision tree for iris species

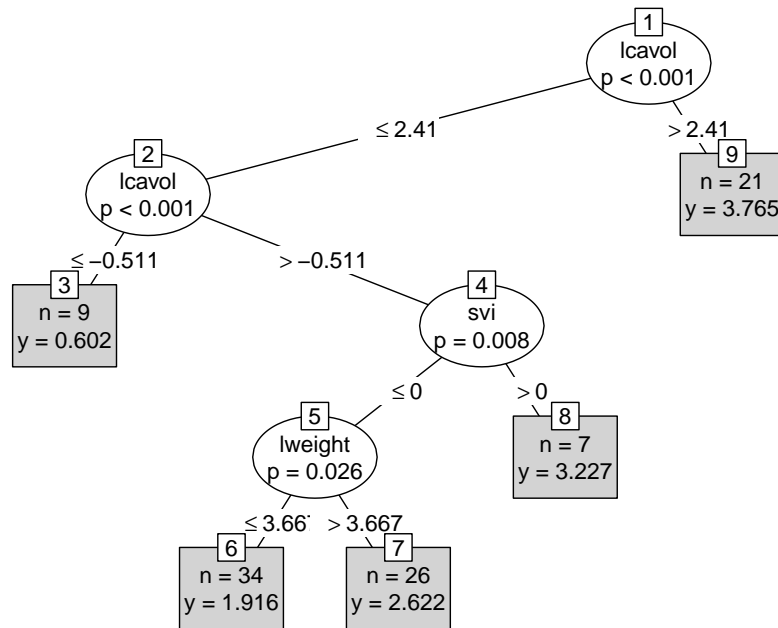


Figure 10.2: Regression tree for log psa in prostate data




Figure 10.3: Plot of samples and computed disparity at each location

are not restricted in that fashion. In particular, they tend to be more flexible than linear models, as they can automatically discover interactions.

10.1.3 Learning trees

Although it is possible in some cases to design a tree manually from expert knowledge, we will usually wish to learn the tree by an automated method. To learn such a tree from a dataset, we need to learn a sequence of dichotomies.

Most common algorithms can be classified as *greedy* algorithms, which attempt to learn the next best dichotomy at every step. To do so, it considers every variables, and attempts to choose the branching point that maximizes the reduction in disparity, which is usually defined as for a cut at c to be:

$$\text{disparity} = \sum_{i:x_i < c} (y_i - \mu_1)^2 + \sum_{i:x_i \geq c} (y_i - \mu_2)^2, \quad (10.1)$$

where μ_1, μ_2 are the group means defined by:

$$\mu_1 = \sum_{i:x_i < c} y_i \text{ and } \mu_2 = \sum_{i:x_i \geq c} y_i. \quad (10.2)$$

This criterion allows the tree to identify good locations to break the sample, and is a good criterion when the true value is piecewise constant, for example as in fig. 10.3. However, just like any greedy procedure, it can be misled if the initial cuts chosen are very suboptimal.

10.1.4 Random forests

Although simple trees can learn complex rules when given enough data, they tend to be somewhat rigid as they have to learn piecewise constant functions. In addition, the greedy tree growing algorithms can be sub-optimal on the particular dataset that is used for training. To reduce these problems and improve the quality of tree-based models, it is common to use an *ensemble* of trees, called *random forest*.

A random forest, as the name implies, is a predictor built off a large number of trees (usually 500 or more), with each tree grown on a somewhat randomised sample of the data, in a fashion similar to the bootstrap (c.f. section 4.5.3). The final prediction is then obtained by averaging the prediction of each individual tree. This improves the estimate by reducing the variance of the predictor.


Random forests are one of the most general and popular methods for prediction today. They can be applied to both categorical and numerical data in a straightforward fashion, and can predict both categorical and numerical outcomes. However, as the prediction is obtained by averaging hundreds of different trees, it is often difficult to understand which variables drive the decision.

10.2 k -nearest neighbours

k -nearest neighbours are a *non-parametric* prediction method, where we do not fit a specific model but instead make use of the entire sample to produce predictions for new observations. The main idea of the k -NN estimator is to consider, for each new observation, the most similar observations in our training sample. From these observations, we may then expect the outcome of the new observation to be similar to that observed in the training set.

More formally, the k -nearest neighbour estimator is defined (for k fixed) by the following procedure: to predict the outcome for an new observation with covariates x , consider the k closest observations x_1, \dots, x_k which minimize the distance between x and x_i in the training set. The prediction is then given by the average value of the observations $n^{-1} \sum_{i=1}^k y_i$. For a categorical outcome, we may instead predict the category that corresponds to the majority of the neighbouring observations.

Here, k is a tuning parameter which describes a bias-variance trade-off. As k becomes larger, the fitted model becomes smoother and less sensitive to noise. On the other hand, as k becomes smaller, the fitted model becomes more flexible but more sensitive to noise. Choosing an appropriate value of k is essential to ensure a good performance of the estimator, and is usually done using cross-validation. For example, we have plotted the classification regions for $k = 5, 15, 60$ in a simple example in fig. 10.4. In this case, we see that with $k = 5$, there is some noise and the region is discontinuous, whereas with $k = 60$, the classification region is very smooth but cannot capture the full complexity of the phenomenon.

Figure 10.4: Example of k -NN classification for varying k

11 Advanced topics in statistics

In this chapter, we will discuss some more advanced topics that mostly explore cases where the simple i.i.d. experimental model is violated, and what we can do in those conditions.

11.1 Time Series

We start our discussion by considering cases where our data is not distributed independently at random. Indeed, in many cases, the observations we obtain might have some correlation, and it is important to account for them.

The most common case of observations that have correlation are time series, which are sequences of observations throughout time. They appear frequently in context such as economics and finance. In such contexts, the value we observe sequentially are correlated. For example, the unemployment rate from last month are somewhat predictive of those next month (hence not totally independent), and our models should account for such patterns.

11.1.1 Stationarity

In addition to correlation in the models, we must also account for trends in the data. Indeed, in order for statistical predictions to be meaningful, it must be that the data we are using to understand the patterns are similar to the data we wish to predict for. In the context of time series, this implies in particular that the future data must be statistically similar to the past data. This condition is called *stationarity*, and is essential to analyzing time series data.

Indeed, when time series fail to be stationary, we may be misled by common trends due to the progress of time. For example, consider the plot of the number of people dying tangled in their bedsheets and that of total cheese consumption in fig. 11.1. The correlation between the two series is measured at 94%, an abnormally high level of correlation. However, this is mainly due to the fact that both quantities increase with time, violating the stationarity assumption and making our statistical statement dubious.

We may instead look at the yearly difference in the time series, which shows much lower correlation at about 33%, which is not surprising given a specifically chosen example. In this case, although the value of the series is not stationary, as it is increasing (future values are in general higher than past values, violating the condition that they are “statistically similar”), the difference between each time point *is* stationary (the yearly increase in the future “looks like” the yearly increase in the past). In such cases, we say that the series is *integrated*, and we should consider working on the differentiated series instead.




Figure 11.1: Total cheese consumption and number of people dying tangled in bedsheets, total value and yearly difference.

11.1.2 Measuring self-correlation

An important feature of time series is that the future values exhibit dependency on past values. In order to measure this dependency, we may define two functions: the *autocorrelation* function, and the *partial autocorrelation* function.

The autocorrelation function measures the correlation of the series with itself at a given lag i :

$$\text{acf}(i) = \text{Cor}(X_t, X_{t-i}). \quad (11.1)$$

That the autocorrelation does not depend on the specific time point t at which it is measured is due to the hypothesis of stationarity.

On the other hand, the partial autocorrelation measures the conditional correlation between a time series and its past. Indeed, if a time series is correlated with itself at lag 1, then we are claiming that X_t and X_{t-1} have some correlation. However, a similar reasoning implies that X_{t-1} and X_{t-2} have some correlation, and by transitivity it is not unlikely that X_t and X_{t-2} have some correlation simply due to that.

The partial autocorrelation function at lag 2, say, measures the additional dependency of X_t on X_{t-2} , once we have accounted for the correlation which flows through X_{t-1} . It can be defined as conditional correlation, namely:

$$\text{pacf}(i) = \text{Cor}(X_t, X_{t-i} \mid X_{t-1}, \dots, X_{t-i+1}). \quad (11.2)$$

11.1.3 Common time series model

In the context of time series, we are interested in modelling explicitly the value of the next observation as a function of the past observation. In this context, the most commonly used model is called the *autoregressive* model, often written $AR(p)$ where $p > 0$ is an integer.

The autoregressive model describes future values as a linear regression on the past values, given by:

$$X_t = \alpha + \beta_1 X_{t-1} + \dots + \beta_p X_{t-p} + \epsilon_t \quad (11.3)$$

where ϵ_t is some i.i.d. noise, often called *innovation*.

The *moving average* model, or $MA(q)$ for $q > 0$ an integer, describes future values as a moving average of latent values, given by:

$$X_t = \theta_1 \epsilon_{t-1} + \dots + \theta_q \epsilon_{t-q}. \quad (11.4)$$

It is common to combine these two models into the $ARMA(p, q)$ model, which is a flexible model that is most commonly used. The ARMA model is able to capture most normal time series behaviour. In practice, the values of p and q are chosen using some model selection algorithm, such as AIC (in a similar fashion as in section 7.3). However, these models are not always very interpretable.




Figure 11.2: Number of international airline passengers 1948-1960

11.1.4 Seasonality

An important feature of time series model is that of seasonality. In many practical applications, time series exhibit a cyclical pattern generally called seasonality. It is not uncommon for models to exhibit yearly, monthly, weekly or daily seasonality. It is important to account for such seasonality when creating models, as it is an important feature of the data.

For example, consider the time series of the number of airline passengers from 1948 to 1960, as plotted in fig. 11.2. We see that the number of passengers follows an increasing trend, reflecting the development of international air travel during that period, but also displays a strong seasonal component.

11.2 Mixed effects

Although our examples of statistical modelling have mostly dealt with i.i.d. observations, a very common occurrence in economic and social science models is that the observations are grouped by units, for example in the context of *panel data* or hierarchical data. In such datasets, some observations may be logically grouped together. For example, consider a sleep deprivation study where each individual is placed under a sleep deprivation condition, and their reaction time measured every day over 10 days. Then, we have a

measure for each day for each subject. However, it is natural to group all the measures corresponding to one subject together as they may have common aspects. Another example could be attempting to measure the level of proficiency in mathematics for a given student. The level of proficiency for each student probably depends on individual aspects of the student, but also the classroom they are placed in – which some students might share – and the school they are placed in, which may regroup classrooms.

11.2.1 ANOVA

Mixed effects models can be seen as a generalization of the problem of *ANOVA*, or analysis of variance. This problem attempts to characterise whether the average outcome of a numerical variable differs across categories. For example, suppose that we collect data on interest rates for new car purchases across different cities. We would like to test the hypothesis the average rate differs among these cities.

A natural operation to do so is to compare the variance of the data among each city, with the total amount of variance observed. Indeed, let y_i be the rate for person i , and $x_i = 1, \dots, l$ the city in which person i lives (numbered from 1 to l). Suppose that we have n total observations, with n_k observations in the city k . Then, it is natural to compare the residual sum of squares:

$$\text{RSS} = \sum_{k=1}^l \sum_{i:x_i=k} (y_i - \mu_k)^2 \quad (11.5)$$

and the total sum of squares:

$$\text{TSS} = \sum_{i=1}^n (y_i - \mu)^2, \quad (11.6)$$

where the μ_k are the group means and μ is the total mean defined by:

$$\mu_k = \frac{1}{n_k} \sum_{i:x_i=k} y_i \text{ and } \mu = \frac{1}{n} \sum_{i=1}^n y_i \quad (11.7)$$

This corresponds to partitioning the total sum of squares, or total variance, into two components: the residual sum of squares, or RSS, and the explained sum of squares, or ESS. The explained sum of squares correspond to the amount of variation in the data that may be explained by partitioning our data. If we can explain a large amount of the variation by partitioning our data, we may be tempted to claim that the partition is informative, and that the different categories have different means. On the other hand, if we can only explain a small amount of the variation, we may be tempted to say that the partition is uninformative.

In fact, ANOVA is equivalent to linear regression with a categorical variable. Indeed, the RSS and TSS correspond to their definition when regressing a numerical outcome on a single categorical outcome, as for a regression on a single categorical variable, we may write the equation as:

$$\mathbb{E} y_i = \mu_k \text{ for } x_i = k, \quad (11.8)$$

where μ_1, \dots, μ_l denote the mean for each group. The test for different means is then equivalent to a total significance test of the entire linear regression.

11.2.2 Mixed effects

However, suppose that we wish not only to understand whether the means are different between each group, but also what the average value for a new group could be. In the example of the interest rates, suppose that we have sampled 6 cities at random from a given state, and wish to predict the average interest rate in another city. Then our analysis above is less helpful, as although we have a value μ_k for each city, we have no information about what value the next μ_{l+1} could be.

Random intercepts In this context, we would like to model μ_k as random, as we are sampling our cities randomly, and are not so much interested in the value for each city, but rather understanding what the average value for each μ_k could be. Thus we can model the average rate for each city μ_k as a normal itself, $\mu_k \sim \mathcal{N}(\mu, \sigma^2)$, with some grand mean μ and variance among cities σ^2 .

This notion of a random coefficient captures the essence of a *random effect*, which describe effects that follow some design not set by the experimenter and not of direct interest to the experimenter. In the context of the example mentioned at the start of the section, this can be the effect of the school, or the individual behaviour of a person when faced with sleep deprivation.

Random slopes These random coefficients can occur both in terms of random intercepts, and random slopes, where the effect of a given variable depends on the subject. For example, as we can see in fig. 11.3, the progress of decrease in reaction time for each subject is somewhat different. In this context, we may model the reaction time y at day x for subject k as:

$$\mathbb{E} y = \alpha_k + \beta_k x, \quad (11.9)$$

where the intercept α_k and the slope β_k depend on the subject. To better understand how human subject react to sleep deprivation, we may be interested in modelling β_k as a random effect, putting $\beta_k \sim \mathcal{N}(\beta, \sigma_\beta^2)$, where β is then the average increase in reaction per day of sleep deprivation, and σ_β^2 the variance of that quantity among human subjects.




Figure 11.3: Reaction time per day of sleep deprivation for each subject

11.3 Missing data

Missing data often arises in practice depending on the fashion in which the data was collected. In science, numerous instruments have a natural failure rate, and may fail to record some small percentage of the data. In social sciences, data is often collected through surveys, and surveyed individuals may often decline to answer certain questions. In these cases, it is important to understand how the missingness can affect statistical conclusions, and how to deal with missing data.

11.3.1 Types of missingness

In order to understand missingness from a statistical perspective, we need to define it as a statistical model. We will model missingness as some random process, and we discuss several possible assumptions on the model.

The most restrictive such assumption is called *MCAR*, or missing completely at random. In this case, whether a given observation or variable is missing is completely random, and does not depend neither on the value of the other variables, nor on the “true” value of the missing variable. Although this assumption can hold in certain cases involving large arrays of instruments that may have some random failures (such as in gene sequencing), this is in general a too restrictive assumption.

The most general assumption is called *MNAR*, or missing not a random. In this context, whether a value is missing may depend on the “true” value. For example, individuals with depression may be less willing to answer a question on whether they have depression, or an instrument to measure wind speed might not function above a certain windspeed. This case is impossible to solve in a generic manner, and requires specific modelling of the missing process – often relying on domain knowledge.

The assumption that is commonly used in statistics is called *MAR*, or missing at random. In this case, whether a value is missing *may not* depend on the “true” value that is hidden, but *may* depend on the other values that are not missing. For example, it may be that men are less likely to answer a question on whether they are depressed, regardless of their actual state of depression. If men are also more likely to be depressed, then this might introduce a bias in underestimating the depression (as men are both more likely not to answer the question, and are more likely to be depressed). However, in this case we may adjust for the bias by controlling for the gender of the subject, and this will be the case of most interest for us.

The MAR at random assumption cannot be verified from the data (as it is a statement about how the *missing* values behave, which we do not observe by definition), it is more believable the more data we have. Indeed, consider the previous example of depression. If we did not record the gender of the subject, depression would no longer be missing at random as it would appear to us that individuals with more depression are less likely to answer the question. As we collect more covariates and adjust for them, it is more likely that the missing at random assumption is a good approximation.

11.3.2 Non-response and response biases

The problem of creating surveys that elicit true answers from surveyed individuals is delicate, and at the intersection of many fields. In some cases, statistical tools may help us reduce the impact of distortions caused by surveying individuals instead of being able to collect the true value.

In this section, we will mostly discuss missing data, which can often appear in surveys due to the so-called *non-response bias*, whereby surveyed individuals decline to answer some questions depending on their characteristics. So long as we may use the data that we have gathered (e.g. gender, socio-economic status, etc.) to predict whether an individual may answer a certain question (e.g. income), we are in the MAR case, and may use the techniques described in this section to handle the data.

On the other hand, surveyed individuals will also often state inaccurate responses when asked, often called the *response bias*. This can happen for a large number of reasons. For example, the phrasing of the question may bias respondents towards answering in a specific fashion. In other cases, the respondent may feel that society places some judgement on the answer of a certain question (most commonly concerning subjects such as smoking, number of sexual partners, etc.), and may feel compelled to give a biased answer, often known as the *social desirability* bias. Finally, when asked to recall some numbers (e.g. how often do they exercise, how many types of fruits have they eaten yesterday etc.), respondents tend to round the numbers, creating some biases towards round numbers (5, 10, 50, 100 etc.). Each of these phenomenon must be modelled separately if we wish to take them into account, although we will not discuss these problems in this section.

11.3.3 Complete case analysis

The simplest method to deal with missing data is the so-called *complete case analysis*, which consists of simply ignoring the observations with missing data. This has the advantage of being an extremely simple method, that can be used in every context. However, complete case analysis suffers from two main problems: it can be biased, and we may lose a substantial proportion of our data.

Sample size Complete case analysis requires us to exclude all samples where we have some missing observations. This can be acceptable when the rate of missingness is small (e.g. less than 5%). However, as the amount of missing data increases, the number of available observation decreases, and we have less data available to perform our analysis. This can be particularly problematic as the number of variables increases. For example, in typical genetic analysis, less than 5% of the gene SNPs are missing for each individual. However, as each individual has on the order of millions of SNP locations, it is pretty much guaranteed that we have no complete record for when given individual, in which case complete case analysis fails.

Biases Complete case analysis is correct when the data is missing completely at random. However, it may be biased when the data is missing at random if no attempt is made to correct for the covariates. For example, suppose that white college educated men tend to be less likely to report their income, and also tend to have higher income. Then, if we were to compute an average income on only the available data, that average may be biased as we have excluded subjects with higher income on average. However, if we have included the race and gender in the model, and formed a regression on the income, we can avoid the bias. In general, complete case analysis is only unbiased in the MAR setting if the analysis adjusts for the covariates causing the missingness.

11.3.4 Multiple imputation

Multiple imputation is a general and popular method to solve missing data problem in the MAR setting. The basic idea of *imputation* is to “fill-in” the missing data by guessing from the observed data. For example, if a white college educated man declines to report his income, we may fill in that value from the income reported by other white college educated men. As long as we are in the MAR setting, the income for the individual that declined to include it is stastically similar to that of similar individuals, and so this produces a reasonable estimate.

By imputing and completing the dataset, we may thus obtain a dataset that has no missing values, and is on average representative of the true data without missingness. We may then use any method of our choice on the completed dataset to analyze the data, and obtain unbiased estimates.

However, although these estimates are unbiased, they tend to be overconfident (i.e. the confidence intervals that we obtain are too narrow). Indeed, in the completed dataset, we have replaced unknown values by fictitious fixed values that we have guessed. However, as we have only done this once, we are now underestimating the variance of the unknown observations. In order to compensate for this fact, we often use *multiple imputation*, where we complete several datasets randomly. We may then combine the estimates from each of those random completed datasets to estimate the uncertainty due to not knowing the exact missing value.

Index

- χ^2 test of independence, 70
- k -NN, 92
- AIC, 74
- alternative hypothesis, 67
- ANOVA, 99
- anscombe, 13
- association, 11
- asymptotic, 58
- autocorrelation, 97
- autoregressive, 97
- bandwidth, 17
- bar chart, 15
- Bayesian, 26
- Bernoulli, 38
- best subsets, 75
- bias-variance, 57
- BIC, 74
- Binomial, 39
- bootstrap, 61
- boxplot, 18
- categorical, 9
- cauchy, 49
- centrality, 11
- circular
 - variable, 9
- classification, 85
- complement, 26
- complete case analysis, 103
- conditional, 13
- confidence interval, 60
- confounding, 51
- conjugate, 80
- consistency, 58
- contingency table, 12
- continuous
 - random variable, 35
- convolution, 45
- correlated, 12
- correlation, 12
- covariance, 11, 48
- cross-validation, 76
- data frame, 8
- data matrix, 8
- decision tree, 89
- density, 35
- disjoint, 26
- dispersion, 11
- distribution
 - function, 36
- double blind, 52
- dummy coding, 71
- ensemble, 92
- epistemic, 31
- estimator, 56
- experiment, 51
- explanatory
 - variable, 8
- exponential, 40
- faceting, 20
- family, 38
- Fisher's exact test, 70
- frequentist, 26
- gam, 66
- graph, 10
- greedy, 91
- histogram, 15

INDEX

- imputation, 104
- inclusion-exclusion, 27
- independence, 27
- independent
 - random variable, 43
 - variable, 8
- innovation, 97
- instrumental variable, 53
- integrated, 95
- interquartile range, 11
- IQR, 11
- joint distribution, 41
- kde, 17
- kernel density estimator, 17
- kurtosis, 59
- level, 9
- likelihood, 55
- Likert, 10
- linear, 46
- link, 65
- log-likelihood, 56
- lognormal, 44
- loss, 83
- MAR, 102
- marginal, 13
- MCAR, 102
- mean, 11
- median, 11
- mle, 57
- MNAR, 102
- model, 55
- moment, 59
- moving average, 97
- mse, 57
- multiple imputation, 104
- mutually exclusive, 26
- nominal, 9
- non-parametric, 63, 92
- non-response bias, 103
- Normal, 40
- normalized, 33
- null hypothesis, 67
- numerical, 9
- observation, 8
- observational, 52
- odds ratio, 73
- ontic, 31
- ordinal, 10
- ordinary least squares, 63
- p-value, 68
- p.d.f., 35
- panel data, 98
- partial autocorrelation, 97
- pie chart, 15
- pivot, 69
- placebo, 51
- Poisson, 39
- posterior, 79
- power, 68
- precision, 86
- prediction interval, 85
- prior, 79
- prospective, 52
- quartile, 11
- random effect, 100
- random forest, 92
- recall, 86
- regression tree, 89
- response bias, 103
- retrospective, 52
- ridge regression, 76
- sample space, 25
- scatter plot, 17
- size, 67
- skewness, 59
- social desirability, 103
- square loss, 83
- standard error, 56
- standard normal, 40
- stationarity, 95

total probability, 30

type I, 67

type II, 68

unbiased, 56

uncorrelated, 48

uniform, 39

variable

data, 8